

DAT650 Blockchain Technology - Course Script

Leander Jehl

November 9, 2022

Contents

1	Data structures	4
1.1	Hash Function	4
1.2	Hash chain	6
1.3	Merkle trees	7
1.3.1	Committing to multiple values	7
1.4	Blockchain	9
2	Transactions	11
2.1	Digital Signatures	11
2.2	Account balances and UTXO	11
2.2.1	UTXO	13
2.2.2	Privacy in the UTXO model	15
3	Proof of work	16
3.1	Centralized straw-man system	16
3.2	PoW function	18
3.3	Block creation with PoW	19
3.3.1	Adjustable difficulty	21
3.3.2	Fees and mining rewards	22
3.3.3	Forks and longest chain rule	23
3.4	PoW and network latency	24
3.5	Attacks on bitcoin mining	25
3.5.1	Stubborn mining	26
3.5.2	51% attack	28
3.5.3	Selfish mining	28
3.6	P2P networking and network layer attacks	29
3.6.1	Inventory messages and delivery denial attack	29
3.6.2	Sybil attack	30
3.7	Attacks on transactions	30
3.7.1	Eclipse attack	31
3.8	Updating a blockchain	32
3.8.1	Soft fork	32
3.8.2	Hard fork	33

3.8.3	Hard and soft forks	33
3.8.4	Analysis	33
3.8.5	Pow as voting	34
4	Alternative PoW	35
4.1	Improving PoW	35
4.1.1	Alternative mining puzzles	35
4.1.2	ASIC resistance	35
4.1.3	Proof of useful work	36
4.1.4	Proof of Storage	36
4.2	Proof of Stake	37
4.2.1	Fairness	38
5	Scaling the Blockchain	39
5.1	Blocksize and block interval	39
5.1.1	Ghost	40
5.1.2	Inclusive chains and uncles	41
5.2	Uncles for scaling	42
5.3	Bitcoin NG	43
5.4	Sharding for blockchains	44
6	System models	46
6.1	Unpermissioned systems	46
6.2	Permissioned systems	46
6.3	Failure models	47
7	A BFT blockchain protocol	50
7.1	Proof of certification	50
7.2	Safety	51
7.3	Liveness	53
7.4	Safety revisited	55
8	Smart Contracts	58
8.1	Ethereum and Smart Contracts	58
8.1.1	SmartContract	58
8.1.2	Accounts	59
8.1.3	Transactions	60
8.1.4	Gas	61
8.2	Smart Contract Security	62
8.3	Oracles	62
8.4	Off chain transactions	63
8.4.1	Uni-directional payment channel	64
8.4.2	Bi-directional payment channel	64
8.4.3	Payment networks	65

9	Hybrid blockchains	66
9.1	Consensus guarantees	66
9.1.1	Bitcoin guarantees	66
9.1.2	BFT guarantees	67
9.2	Hybrid BFT	67
9.2.1	ByzCoin	68
9.2.2	Casper FFG	69
9.3	Algorand	70
10	Use Cases	72
10.1	Supply chain	73
10.2	Inter bank payments	73
10.3	Digital services	73

Chapter 1

Data structures

1.1 Hash Function

Definition 1. A (cryptographic) **hash function** H ($H(x) \mapsto y$) is a *deterministic* function that maps inputs (x) of arbitrary size to "small" outputs (y), e.g. 256 bit numbers.

Note 1. The idea for a hash function is that the result should "look" random. However the function is *deterministic*, i.e. the same input always results in the same output.

Note 2. A mental image for a hash function is the following:

Imagine a gnome that lives in a closed box. When you give a new piece of text to the gnome, he writes down the text, rolls some dice and gives you a random number for your text.

When given the next piece of text, the gnome will look up if he has seen this text before. If yes, he will return the number from before. If no, he rolls the dice and gives you a new number.

Example 1. *Example for hash functions are*

MD5 *not secure*

RIPEMD-160 *outputs 160 bits - not secure*

SHA-1 *not secure*

SHA-2 *actually SHA-256 and SHA-512 give 256 and 512 bit output.*

SHA-3 *standard from 2015*

On unix terminal (also mac) compute SHA-256 of a file

```
shasum -a 256 file.txt
```

or

```
printf "hello world" | shasum -a 256
```

Hashes here are given in hexadecimal numbers.

Security summary at

https://en.wikipedia.org/wiki/Hash_function_security_summary

Definition 2. For a secure hash function the following security properties need to hold:

Pre-image resistance (one way) Given a d -bit y it is infeasible to find any x such that $y = H(x)$

Weak collision resistance Given an input x it is infeasible to find a different x' such that $H(x) = H(x')$.

Collision resistance It is infeasible to find inputs x and x' such that $x \neq x'$ and $H(x) = H(x')$.

Infeasibility means that an exhaustive search is the best strategy to find such values and the probability to find an x as above is smaller than $\frac{1}{2^{112}}$

Example 2. *Examples for the use of cryptographic hash functions are:*

- *Including hashes of third party content in HTML-tags, e.g. scripts.*
- *Not storing user passwords in plain text.*

Note 3. a) Given a secure hash function H that takes strings as input, we can create a hash function H' that takes lists, structs or byte arrays as input.

b) For pre-image resistance to be useful, the domain of structs should be large and structs should be unpredictable.

c) Given a secure hash function H we can create a family of hash functions H_i by concatenating a number or letter i to the input of H $H_i(x) = H(x||i)$. We can choose i at random.

Note 4. Given a hash function $H : \mathcal{X} \rightarrow \mathcal{Y}$ that hashes strings (or byte slices) we can create a hash function $H : \mathcal{X} \times \mathcal{X} \rightarrow \mathcal{Y}$ by concatenating the inputs.

We write $x_1||x_2$ for the concatenation of inputs x_1 and x_2 . Note that the resulting hash function is not commutative.

$$H(x_1||x_2) \neq H(x_2||x_1)$$

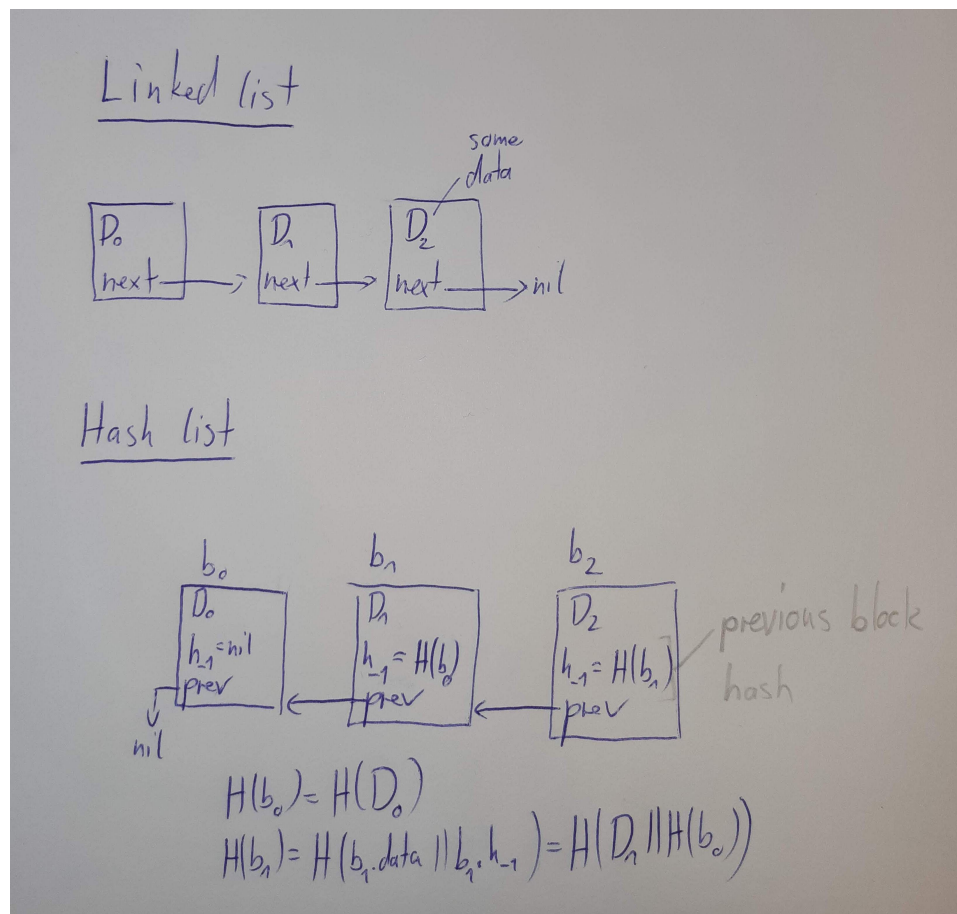
1.2 Hash chain

Assume that H is a secure hash function that maps a struct to a 256-bit.

Definition 3. A **hash-chain** is a data-structure composed of nodes called blocks. Every block b contains a pointer $b.prev$ to a preceding block and a field $b.h_{-1}$ that contains the hash of the block referenced by $b.prev$. A block also contains a data item $b.data$.

- The hash of a block b is computed as $H(b) = H(b.h_{-1} || b.data)$.
- $b.h_{-1} = H(b.prev)$
- There exist one block b_0 , for which $b.prev$ is empty.
- For two blocks b and b' , if $b.prev = b'$ we say that b is a *successor* of b' .

In a **hash-chain** every block has at most one successor.



Lemma 1. *The blocks in a hash-chain can be linearly ordered according to the successor relation with the block b_0 as the first block.*

Note 5. If blocks b_0 , b_1 and b_2 contain data D_0 , D_1 and D_2 then

$$\begin{aligned} H(b_0) &= H(D_0) \\ H(b_1) &= H(H(b_0)||D_1) = H(H(D_0)||D_1) \\ H(b_2) &= H(H(b_1)||D_2) = H(H(H(b_0)||D_1)||D_2) \end{aligned}$$

Example 3. *Assume a notary maintains, a hash-chain of all documents it has signed.*

- *This allows document holders to proof that their document was signed.*
- *This allows to document holders to proof which document was signed first.*
- *Prevent the notary itself to change in which order documents were signed or remove single documents.*

1.3 Merkle trees

By publishing a hash of a document/data we can commit to this document, without revealing it. In the following we look at how to publish many such hashes concurrently.

Example 4. *Given a document describing an invention, I can publish the hash of the document in a newspaper. If someone else then tries to claim my invention, I can proof that I had the document before the date in the newspaper.*

This does not reveal my invention.

Example 5. *In a election you can hash your vote and publish the hash of your vote. After some deadline, you can then reveal your vote to be counted.*

Problem: *An adversary can simply hash all possible values to vote for and thus discover your vote.*

Solution: *Hash vote, concatenated with random value.*

1.3.1 Committing to multiple values

Given data D_1 , D_2 , D_3 and D_4 . We want to commit to all four values. We can either publish individual hashes or a single hash for all:

- a) Publish $[H(D_1), H(D_2), H(D_3), H(D_4)]$.
- b) Publish $H(D_1||D_2||D_3||D_4)$.

Variant b) requires to publish only a single hash. However to proof that D_1 was committed to, it is necessary to reveal also D_2 , D_3 , and D_4 .

Example 6. We can build on scheme a) and additionally create $h_{1,2} = H(H(D_1)||H(D_2))$, $h_{3,4} = H(H(D_3)||H(D_4))$ and $h_{1,2,3,4} = H(h_{1,2}||h_{3,4})$.

We publish $h_{1,2,3,4}$.

To prove that D_3 was included we present D_3 , $h_4 = H(D_4)$ and $h_{1,2}$. To check the proof, recompute $h_{3,4}$ and $h_{1,2,3,4}$.

Definition 4. A **Merkle Tree** is a (binary) tree where every node is labeled with a hash h .

- For internal nodes h is the hash of the concatenated labels of its children ($h = H(\text{leftchild}.h||\text{rightchild}.h)$).
- For a leaf node, h is the hash of the data stored in the node ($h = H(D_i)$).




Figure 1.1: Merkle tree with root $h_{1,2,3,4}$

Lemma 2. If h_r is the root of a Merkle Tree with n data elements, it is possible to proof that D_i is one of the data elements by revealing D_i and $\log(n)$ node labels, i.e., hashes.

Example 7. In a Merkle Tree with 16 data elements, an inclusion proof contains 4 hashes.

In a Merkle Tree with 1000 data elements, an inclusion proof contains 10 hashes.




Figure 1.2: Hashes necessary and computed to proof inclusion for D_3 .

In a Merkle Tree with 1 million data elements, an inclusion proof contains 20 hashes.

Note 6. It is possible to enhance other tree data-structures, e.g. binary search tree or radix tree with hashes from a Merkle tree.

1.4 Blockchain

The following blockchain definition combines hash-chain with Merkle trees.

Definition 5. A **Blockchain** is a hash-chain where the data entry in every node is the root of a Merkle tree.

Note 7. The blockchain construction was not invented for bitcoin. It is used previously in *linked timestamping*, where block headers are published in newspapers.




Figure 1.3: A binary search tree enhanced with hash references.

Chapter 2

Transactions

2.1 Digital Signatures

A short recap of digital signatures.

A signature scheme comprises key generation, signing and verification functions. Key generation creates a public and secret (private) key pair (pk, sk) .

Signing function takes an arbitrary message m and the secret key sk and outputs a signature σ .

$$\sigma \leftarrow \text{Sign}(m, sk)$$

A verify function takes a message m , signature σ and the public key pk . The verification function returns **true**, if σ was produced for m with the secret key sk matching pk . Otherwise verification returns **false**.

$$\{\mathbf{true}, \mathbf{false}\} \leftarrow \text{Verify}(m, \sigma, pk)$$

Bitcoin currently uses ECDSA signatures, i.e. signatures where the public key is a point on an elliptic curve.

Example 8. *One example for digital signatures is https certificates and the establishment of a TLS session. These certify that a given public key belongs to a certain domain. They are signed by a Certificate authority. The public key of the certificate authority is typically already included in your browser installation.*

The binding of a public key to a specific entity (e.g. web-domain) is called public key infrastructure (PKI).

2.2 Account balances and UTXO

The most prominent use case for blockchains is currently digital currency. We now spend some time to look at how to implement money transfer.

Idea 1. *We can use public keys as identity for users or owners of money.*

Note 8. Idea 1 allows to avoid the use of *public key infrastructure* (PKI), that is otherwise used to bind public keys to identities.

Algorithm 1 shows a simple bank that maintains balances. For every public-key, the bank stores one balance. The algorithm has two transactions, CREATE creates new balances. TRANSFER allows to transfer money from one account (*pk-from*) to another (*pk-to*). The algorithm assumes that balances that have not been initialized have the value 0.

Authentication of transactions poses a problem:

- For CREATE transactions we assume that they are only valid during setup of the system.
- For TRANSFER transactions it is desired that all transactions can be submitted by users. To authenticate the sender, the transaction includes a signature σ . In Line 5 we therefore check that the sender *pk-from* signed the transaction.

A bank based on Algorithm 1 may be susceptible to **replay attacks**. A TRANSFER transaction signed by *pk-from* may be submitted multiple times, given that the account of *pk-from* has sufficient funds. This will result in additional funds transferred to *pk-to*.

Later, we will see, how replay attacks in this account based model can be avoided using a counter/sequence number for each account.

Algorithm 1 Simple Bank using account balances

```

1: balances := [pk]uint
2: CREATE(value, pk)
3:   balances[pk]+ = value
4: TRANSFER(value, pk-from, pk-to,  $\sigma$ )
5:   if verify(m, pk-from,  $\sigma$ ) then           ▷ m = value||pk-from||pk-to
6:     if balances[pk-from] > value then
7:       balances[pk-from]- = value
8:       balances[pk-to]+ = value

```

Note 9. + Transfer transaction can be submitted by a sender without the receiver being online.

- It is possible to (accidentally) send money to a public key that does not exist, i.e. that nobody knows the private key for.

Who runs the bank? In a centralized system, a trusted party could process transactions, compute balances and distribute balances to all parties.

Without a trusted party, it is necessary that transactions are distributed to everybody. Every party can then deterministically compute the balances.

If all participants get all transactions, they can each process them individually and arrive at the same balances.

2.2.1 UTXO

Bitcoin does not use balances. Instead it uses the *unspent transaction output* (UTXO) model. We now explain this model:

Idea 2 (UTXO). *In the UTXO model, instead of storing a balance for each private key, we store for every coin, which private key it belongs to. This is complicated a bit, since we want to be able to split and merge coins.*

Therefore, instead of individual coins, we store unspent transaction outputs, i.e. money received and not yet spend.

Definition 6. A **transaction output** is a tuple (v, pk) that shows, that v funds have been transferred. pk is a *spending condition* that must be met to spend claim v . Typically pk requires a signature with a given public key.

A **transaction input** is a tuple consisting of a reference to a transaction output and an argument that meets the outputs condition. I.e. $(outp_i, \sigma)$ where σ is *redeeming argument* a matching pk , e.g. a signature.

A **transaction** is a tuple containing a list of transaction inputs and a list of new outputs.

Note 10. In bitcoin transactions are implemented in the following way:

- An output from transaction t is identified by a tuple (h_t, i) , where h_t is the hash of t and i is the index in the list of outputs in t .
- Algorithm 2 shows how a transaction is validated. For a transaction to be valid, *all inputs must be unspent*, *input signatures must validate* and the *sum of input values must be larger than the output values*.
- Algorithm 2 ensures that a transaction can only be validated once and no two valid transactions can spend the same output.
- The difference between transaction inputs and outputs is called *transaction fee*.
- Example 10 gives an example for more complex conditions that may be required to spend an output.
- When the value of inputs is larger than the desired value to be spend, it is common to create an additional output that contains change.

Algorithm 2 Transaction validation and maintenance of UTXO

```
UTXO := map[(h, i)] → (value, pk)
TRANSFER(inputs, outputs)           ▷ Transaction t with hash h_t
  for ((h, i), σ) ∈ inputs do
    if UTXO[(h, i)] does not exist then
      abort                           ▷ invalid transaction
    if verify(h_t, σ, UTXO[(h, i)].pk) == false then
      abort                           ▷ invalid transaction
  if sum of values of inputs < sum of values of new outputs then
    abort                             ▷ invalid transaction
  for ((h, i), σ) ∈ inputs do
    UTXO[(h, i)] = nil                ▷ output spent
  h_t := hash(transaction)
  UTXO[h_t] = outputs                 ▷ add new output
```

Example 9. Alice’s public key has two unspent outputs with value 1\$ and 1.5\$. Alice wants to send 2\$ to Bob. To do that, Alice can create a transaction that has her two outputs as input and creates two outputs, one with value 2\$ and Bobs public key. One with value 0.5\$ and Alices public key.

Definition 7. A *double-spend* is a situation where multiple transactions attempt to spend the same output.

Note 11. Note that according to Algorithm 2 only one of double-spend transactions can be validated.

Example 10. The following are the most common examples for arguments necessary to claim a transaction output. In Bitcoin they are expressed in a stack based scripting language.

- a) A signature that matches a certain public key.
- b) A public key that hashes to a certain value and a signature that matches this key. (Pay to public key hash or P2PKH).
- c) Multiple signatures that match a sequence of public keys. (Multisig)
- d) m signatures that match m out of n provided public keys. E.g. 2 signatures from 2 out of 3 specified public keys.
- e) A script that hashes to a certain value and an argument that causes this script to evaluate to true. (pay to script hash)

See *Mastering Bitcoin* book, Chapter 6, *Script Construction* and *Book: Bitcoin and Cryptocurrency Technologies*, Chapter 3.2 for explanation of P2PKH script.

Note 12. To maintain a copy of our transaction based bank, a node has to maintain the set of all unspent transaction outputs *UTXO*.

If variant b) is used instead of variant a) from Example 10 this may significantly reduce the size of the *UTXO* data structure. The same holds, if d) is used instead of c).

P2PKH makes it possible to pay to the hash of a public key. This gives rise to the concept of an address.

Definition 8. An **address** is either a public key or the hash of a public key. Given the address of a user, it is possible to transfer funds to this user, i.e. create an output that this user can claim by providing a correct signature.

Note 13. Bitcoin and many other cryptocurrencies use Base-58 encoding. This encoding uses small and large letters (a to z) and numbers, omitting 0 (number zero), O (capital o), l (lower L) and I (capital i) because of their ambiguity.

Bitcoin addresses use Base58Check encoding which adds a 4 byte checksum before Base-58 encoding, to protect against typos, ...

2.2.2 Privacy in the UTXO model

Different from the account and balance system, the *UTXO* model encourages the use of different addresses. This makes it harder to identify all transactions done by a single user.

However research has shown, that based on transactions, it is easy to identify different addresses belonging to the same user.

On the other hand, *UTXO* allows to trace in which transactions a particular value or coin was involved.

Definition 9. A **tainted coin** is a transaction output that is either the result of a transaction considered unethical or illegal or derived from the output of such a transaction by multiple other transactions.

Note 14. Based on the concept of tainted coins it is debatable whether digital cash based on the *UTXO* model is *fungible*. In economics fungibility is defined as the property that any two units of a good are interchangeable.

The *UTXO* model allows to create a mixing service:

Definition 10. A *mixing service* can be used to prevent a third party from tracking a specific users transactions. A mixing service would receive payments from many users, and pay them back using new addresses.

Note 15. • A mixing service makes it hard to see which of the new addresses belongs to which of the users that sent money to the mixing service.

- A mixing service usually requires a high fee.
- A mixing service is usually implemented as a centralized entity.

Chapter 3

Proof of work

In this chapter we discuss how the Transactions, introduced in Chapter 2 are included in a block in the hash-chain, introduced in Chapter 1.

In Chapter 2 we said that all transactions are broadcast, validated and applied by every node. However applying transactions in different order may cause the system to diverge.

3.1 Centralized straw-man system

We present a straw man solution that relies on a centralized leader to include transactions into a block and issue those blocks. After discussing difficulties with this approach, we present the proof of work based scheme used in bitcoin.

As shown in Figure 3.2, We assume a single leader exists. Transactions are submitted to the leader. At regular intervals, the leader includes all new (and valid) transactions in a merkle tree. The root of the merkle-tree is included in a new block.

The block is then broadcast to all participating nodes, who validate it. *The validation of blocks prevents the leader from including malformed transactions.* However this system still opens several ways for the leader to misbehave:

- A) The leader can omit certain transactions on purpose. (Censorship)
- B) The leader is a single point of failure.
- C) The leader could send different blocks to different processes.

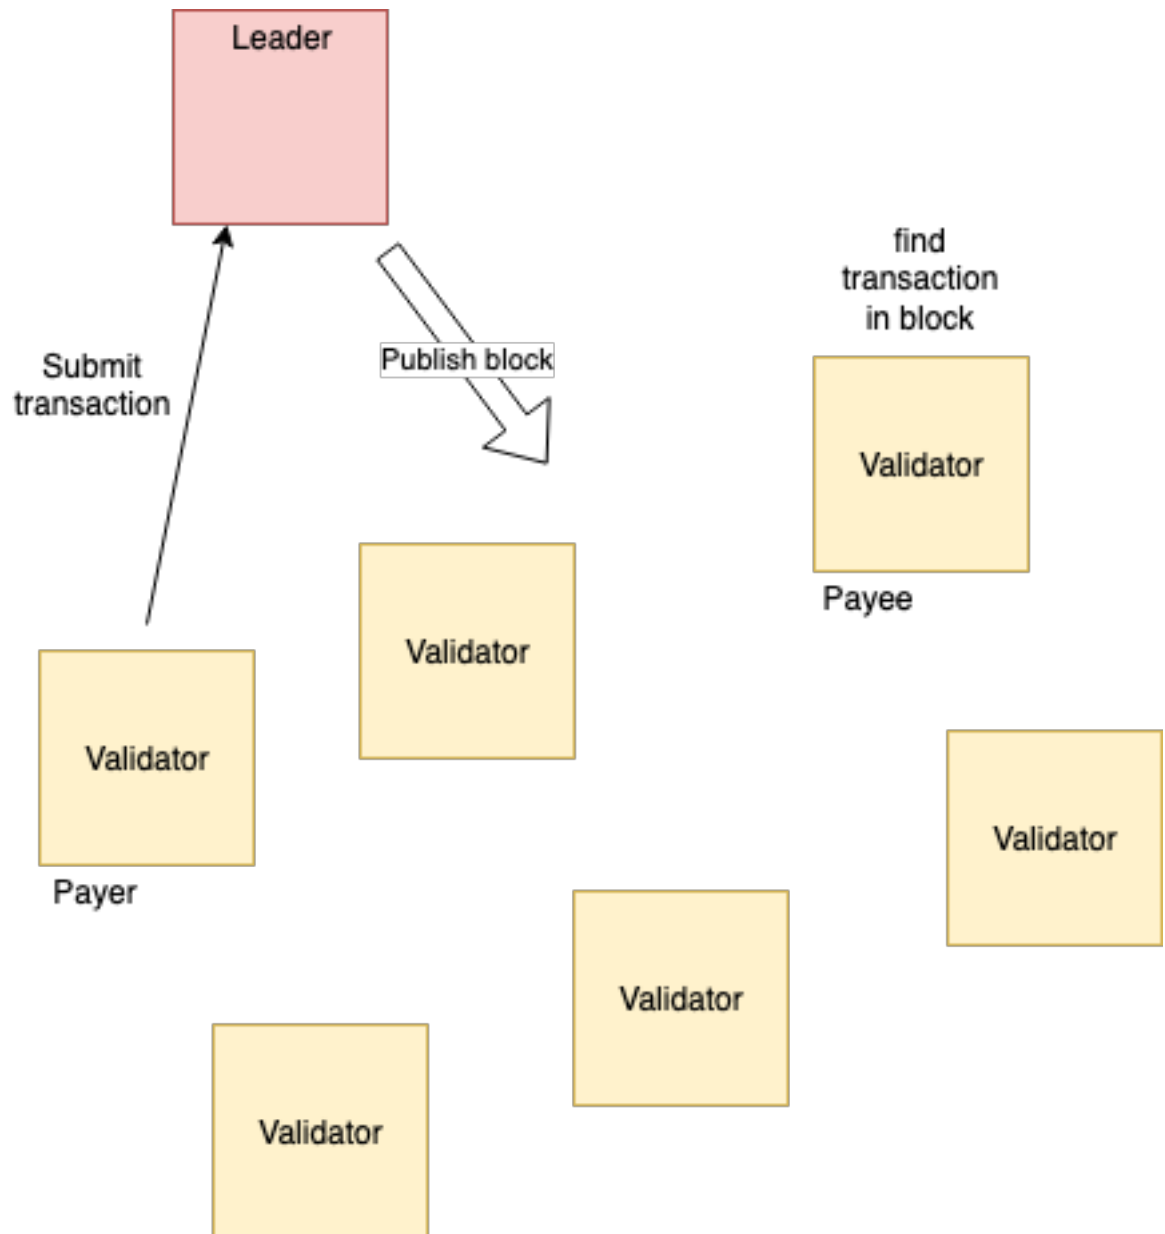


Figure 3.1: Straw-man system with leader. Payer submits transaction to the leader. The payee will eventually find the transaction in a block and know he was payed.

3.2 PoW function

Before we can discuss, how bitcoin avoids a single leader, we have to learn about PoW functions.

Definition 11. (PoW function version 1) For an integer d , the proof-of-work (PoW) function with difficulty d takes a data item and returns a nonce (random bits) and a hash value:

$$(h_{PoW}, nonce) = f_{PoW}(Data)$$

The proof of work is *valid*, if a) h_{PoW} is the hash of the data, concatenated with the nonce

$$h_{PoW} \stackrel{?}{=} H(Data||nonce)$$

and b) the first d bits of h_{PoW} are 0.

Note 16. The function f_{PoW} is computed by choosing a random *nonce* and checking if the condition b) holds. If it does, we also say that the *nonce* *solves* the proof of work.

Lemma 3 below is an important conclusion. It follows since the result of hashing one data item is independent from hashing another data item.

Lemma 3. *Given two nonces, chosen at random. The probability that any of them solves the proof of work is independent.*

Note 17. Lemma 3 has important implications. Especially, it follows that if you already know that many nonces do not solve PoW, you do not have a better chance for the next nonce to be successful.

Theorem 1. *If we conduct multiple, independent Bernoulli trials with success probability p , the expected number of trials necessary until success is $\frac{1}{p}$.*

For proof see here.

Note 18. Using version 1 of the proof of work function, adjusting d the difficulty of the proof of work can only be doubled or halved.

- If $d = 1$, on average, one out of 2 nonces gives a solution.
- If $d = 2$, on average, one out of 4 nonces gives a solution.
- If $d = 3$, on average, one out of $8 = 2^3$ nonces gives a solution.

Definition 12. (PoW function version 2) For an hexadecimal number d , the proof-of-work function with difficulty d takes a data item and returns a nonce (random bits) and a hash value:

$$(h_{PoW}, nonce) = f_{PoW}(Data)$$

The proof of work is *valid*, if a) h_{PoW} is the hash of the data, concatenated with the nonce

$$h_{PoW} \stackrel{?}{=} H(Data||nonce))$$

and b) h_{PoW} written as hexadecimal number is smaller than d .

$$h_{PoW} < d$$

Note 19. The proof of work function from Definition 12 allows to carefully adjust the difficulty d to achieve a certain expected time.

- If $d = 80000...$ with 31 zeros, then one out of 2 nonces gives a solution.
- If $d = 60000...$ with 31 zeros, then one out of 3 nonces gives a solution.

3.3 Block creation with PoW

Idea 3. *The idea behind block creation using PoW is to require that any block published includes a nonce, s.t. the block hash and nonce are results of a PoW function.*

We can thus allow every node to publish a block as outlined below.

- *Transactions are broadcast to all nodes.*
- *All nodes (miners) try to create a new block including new valid transactions.*
- *A newly created block is broadcast to all nodes.*
- *Nodes verify the new block, add it to their state and then start trying to create the next one.*

Definition 13. A *proof of work blockchain* is a blockchain as in Definition 5 where additionally, every block contains a *nonce* and the hash of the previous block hash h_{-1} , the root of the merkle tree and the *nonce* solves the difficulty.

This design give the following properties:

Censorship An adversary cannot prevent a node, it does not control from including a certain transaction.

Failures Failure of individual nodes will not prevent the system from running.

Frequency A high PoW difficulty ensures that blocks cannot be created faster than they can be distributed to members in the network.

Conflicting blocks For a single entity to create two different blocks requires to solve PoW twice. This makes it difficult to publish two conflicting blocks. A high difficulty decreases the probability that two blocks are found concurrently.




Figure 3.2: Block creation with PoW

3.3.1 Adjustable difficulty

Back in 2010 it was possible to join Bitcoin and create a new block using a simple desktop computer. Today specialized hardware (ASICs) is used to solve the PoW.

Idea 4. *To adjust to nodes joining and leaving the system, and nodes using different infrastructure (i.e. hardware) it is possible to adjust the hardness of the blockchain.*

Definition 14. Additional to the merkle root and nonce Bitcoin includes a timestamp in the block (and in the PoW). This allows to compute the average time between two blocks every 2016 blocks and adjust the difficulty.

Note 20. on timestamps in a PoW blockchain

- The time between blocks can vary significantly but this variance has little effect on the average time, taken over 2016 blocks.
- The timestamp can be set by the nodes when creating the block. When validating the block nodes check that the new block has a timestamp within 2 hours of their local clock.

Lemma 4. *If we assume that the probability that the nodes find a new block within time Δ is constant and independent, and assume that the mean time to block creation is 10 minutes, we can compute the probability, that a block is created within t seconds as*

$$P[\text{block created within } t \text{ seconds}] = 1 - \left(\frac{599}{600}\right)^t$$

$$P[\text{no block created within } t \text{ seconds}] = \left(\frac{599}{600}\right)^t$$

Proof. Let p be the probability that a block is found within 1 second. Let T be the random variable describing after how many seconds a new block is found. Since finding a block in a specific second is independent we get:

$$P[T = t] = (1 - p)^{t-1}p$$

From Theorem 1 it follows that $E[T] = \frac{1}{p}$. If we assume $E[T] = 10\text{min} = 600$ we get $p = \frac{1}{600}$. □

Lemma 5. *Assume the current average block interval is Δ_1 . To adjust the average block interval to Δ_2 , we need to adjust difficulty from d_1 to*

$$d_2 := \frac{\Delta_1}{\Delta_2} \cdot d_1$$

Example 11. *If the measured average delay of the last blocks is 5 minutes and the desired delay is 10 minutes you have to make PoW more difficulty (2 times). This is done by decreasing the target value d by 0.5.*

$$d_2 := \frac{d_1}{2}$$

Example 12. *If the measured average delay of the last blocks is 6 minutes and the desired delay is 10 minutes you have to adjust difficulty by*

$$d_2 := \frac{6}{10}d_1$$

3.3.2 Fees and mining rewards

Fees and mining rewards give an incentive to solve the PoW function.

Definition 15. When the sum of outputs of a transaction is larger than the sum of inputs, the difference is called *fee*.

Definition 16. Every block in bitcoin contains a *coinbase transaction*. This transaction has no input and a single output. The output value is the sum of a fixed reward (*mining reward*) and the sum of all fees of transactions included in the block.

Note 21. There are different **pro's and con's for fixed rewards and fees**:

- In bitcoin the mining reward is halved every 4 years. Thus, only a finite amount of bitcoin will ever be created.
- The mining reward allows cheap transactions, since fees do not need to cover full mining expenses.
- Some research suggests, that if mining rewards are small compared to fees, mining becomes unstable.¹
 - Miners may fight over transactions with big fees.
- Mining awards are necessary to get money into circulation.
- Currently in bitcoin mining rewards are 12.5 bitcoin, while fees per block are less than 1 bitcoin.

Note 22. How big should the fee be? Fees are determined by market economics, i.e. nodes choose which transactions to include. Usually those with highest fee.

Processing of transactions requires nodes to use network and processing capacities (for relaying and validating transactions). These expenses depend on the size of a transaction, but not on the value transferred. Therefore:

¹<https://freedom-to-tinker.com/2016/10/21/bitcoin-is-unstable-without-the-block-reward/>

- Fees in bitcoin are usually independent of the value of a transaction. Instead they depend on the size (in bytes) of the transaction.

3.3.3 Forks and longest chain rule

As mentioned above, PoW can reduce the probability for conflicting blocks to be proposed, but cannot prevent this from happening.

Definition 17. A *fork* in a blockchain is when multiple blocks are proposed with the same predecessor. Note that every block in a fork may be extended to a different chain.

A fork is problematic since the different blocks or chains represent two versions of the world.

Imagine your bank being undecided, whether you did or did not pay your rent.

Definition 18. (Longest chain rule) If a fork exists all nodes should adopt the longest chain.

Note 23. While the Longest chain rule is implemented in the standard bitcoin release, it is not enforced. It is possible to extend a different chain than the one given by the longest chain rule. Doing this may change what is the longest chain.

Lemma 6. *Given two chains c_1 and c_2 . To find a new block extending c_1 is equally likely as finding a new block extending c_2 . This even holds, if nodes already spend significant resources to find a block extending c_1 .*

Proof. This follows from Lemma 3. □

Lemma 7. *Assume two chains c_1 and c_2 where c_1 is longer than c_2 . Further assume that most nodes follow the longest chain rule, trying to extend the longest chain. A new block is more likely to be eventually part of the longest chain, if it is added to c_1 rather than c_2 .*

Proof. Chain c_1 is already longer than c_2 . A new block would make c_1 even longer, while it would make c_2 at most equally long to c_1 . □

Example 13. *Figure 3.3 shows a fork. Two processes trying to extend block b_2 found a new block, b_3 and b'_3 .*

Later another block b_4 was found, that extends b_3 . According to the longest chain rule, all nodes will adopt the blocks b_3 and b_4 and discard the block b'_3 .

Transactions that were only included in b'_3 but not in b_3 will not be part of the longest chain. They are effectively discarded.




Figure 3.3: A fork in the blockchain. Two blocks extending b_2 where found.

Coinbase transaction *The coinbase transaction in b'_3 is discarded. Thus the node that created b'_3 loses his block reward.*

Double spending *Assume transactions t and t' spend the same output. Assume transaction t' is included in b'_3 and t is included in b_3 . t' is discarded together with b'_3 . Since b_3 includes t , t' cannot be included in a later block either.*

3.4 PoW and network latency

We now analyze the probability that a fork occurs based on the network delay, i.e. the time it takes to broadcast a block to the different nodes. Results in this section are taken from Decker and Wattenhofer, 2013.

Definition 19. We write δ for the average time it takes for a block until it is validated by a specific node in the network.

Note 24. Based on empirical evaluations, in the current bitcoin network, $\delta = 12.6$ seconds.

In the following theorem we assume that nodes have the same mining power and are following the longest chain rule.

Theorem 2. *Let p be the probability that a block is found within one second. Then the probability for a fork is*

$$1 - (1 - p)^\delta$$

Proof. The probability for a fork is the probability that while the block is propagated, another block is found. Thus

$$P[\text{fork}] = 1 - P[\text{no block found during dissemination}]$$

$$P[\text{no block found during dissemination}] = (1 - p)^\delta$$

□

Corrolary 1. Let $P[\text{fork}^l]$ be the probability that after a fork, both chains are extended by l blocks. It holds

$$P[\text{fork}^l] \leq \left(1 - (1 - p)^\delta\right)^l$$

Proof. Assume one chain c_1 is extended by one block. If a fork of this chain, c_2 is also extended, this has to happen before the new block on c_1 is propagated. Thus the probability that two chains in a fork are extended, is equal to the probability that a block on the second fork is found, while the block on the first fork is disseminated. This is less than the probability of a fork. □

Corollary 1 says that the probability that a block is discarded because of a fork drops exponentially, with the numbers of blocks that are added after this block.

Definition 20. We say that a transaction is *confirmed*, if it is included in a block, and several blocks are added after this block.

In bitcoin, it is recommended to wait for additional 5 blocks, after the block including a transaction.

3.5 Attacks on bitcoin mining

In Section 3.4 we have shown that continued forks are extremely unlikely, if nodes stick to the longest chain rule and do not disturb the network latency.

In the following we first look at two ways to deviate from the longest chain rule, stubborn mining and selfish (hidden) mining. And we look at network attacks that might be performed.

We first look at those attacks from the perspective of fair mining. Then we look at them from the perspective of a double spend.

Definition 21. Mining, or block creation is *fair* if a node that possesses α percent of the hashing power in the network ends up publishing α blocks in the longest chain.




Figure 3.4: Stubborn mining states and transitions.

3.5.1 Stubborn mining

A node does perform stubborn mining, if it does not abandon the current chain for the longest chain. Thus it does not follow the longest chain rule.

More precisely, if there exist two blockchains c_1 and c_2 and c_2 contains one more block than c_1 . Thus the initial state is as shown in Figure 3.3. A stubborn node that has published a block in c_1 that is not part of c_2 will continue to try and extend c_1 until either c_1 becomes the longest. We assume that if the difference between c_1 and c_2 increases, the stubborn node will abandon c_1 .

Theorem 3. *Stubborn mining does not increase the expected outcome of a node, if the node controls less than $\alpha = 0.42$ of the hashing power in the network.*

Proof. We model this system as a markov chain with 4 states.

Loose where chain c_1 was extended faster than c_2 .

-1 where chain c_1 is one block longer than chain c_2 .

0 where c_1 and c_2 are equally long.

Win where chain c_2 became the longest chain.

We ignore the probability that additional forks occur on c_2 but note that they would increase the profitability of stubborn mining. Assume the probability that the attacker finds a block is α , while $\beta = 1 - \alpha$ is the probability that the remaining miners find a block. The states and the transition probabilities are shown in Figure 3.4.

We now calculate the expected number of blocks the Attacker receives with and without doing the attack. For the attack we consider the following cases:

- With probability β the attacker loses in the first step and receives no blocks.
- With probability $\alpha \cdot \alpha$ the attacker mines two blocks. He receive 3 blocks in total.
- With probability $\alpha \cdot \beta \cdot \alpha \cdot \alpha$ the process goes through states $-1 \mapsto 0 \mapsto -1 \mapsto 0 \mapsto \text{win}$. In this case the attacker gets 4 blocks.

Extending the above cases, and omitting those that give 0 blocks, E_{Attack} the expected number of blocks is:

$$E_{Attack} = \sum_{i=0} (i+3) \alpha^{i+2} \beta^i \quad (3.1)$$

$$= 3\alpha^2 + \alpha\beta \left(\sum_i 0(i+3) \alpha^{i+2} \beta^i + \sum_{i=0} \alpha^{i+2} \beta^i \right) \quad (3.2)$$

$$= 3\alpha^2 + \alpha\beta \left(E_{Attack} + \frac{\alpha^2}{1 - \alpha\beta} \right) \quad (3.3)$$

In step (3.2) we used the formula for a geometric sum. Solving the above for E_{Attack} gives

$$E_{Attack} = (3 + \alpha\beta) \frac{\alpha^2}{1 - \alpha\beta}$$

To compute the average number of blocks, the attacker would receive if it did not follow the attack, we note that the he gets one block every time an edge with probability α is traversed. We get the following cases.

- With probability $\alpha \cdot \alpha$ the node mines two blocks.
- With probability $\alpha \cdot \beta \cdot \beta$ the process goes through states $-1 \mapsto 0 \mapsto -1 \mapsto \text{loose}$. The node mines 1 block.
- With probability $\alpha \cdot \beta \cdot \alpha \cdot \alpha$ the process goes through states $-1 \mapsto 0 \mapsto -1 \mapsto 0 \mapsto \text{win}$. In this case the node gets 3 blocks.

Continuing the above analysis, we see that the average number of blocks received when not following the above state machine, but not following the attack, is:

$$E_{NoAttack} = \sum_{i=0} (i+2) \alpha^{i+2} \beta^i + \sum_{i=0} i \beta^{i+1} \alpha^i$$

Using the same techniques as for E_{Attack} , we get:

$$E_{NoAttack} = (2 + \alpha\beta) \frac{\alpha^2}{1 - \alpha\beta} + (1 + \alpha\beta)$$

Plotting both graphs we see that $E_{Attack} < E_{NoAttack}$ holds approximately for $\alpha < 0.42$. \square




Figure 3.5: A 51% attack.

3.5.2 51% attack

If a miner owns $\alpha = 51\%$ of the hashing power in the network, he can attack the network by creating and growing his private chain. Key to this attack is that the attacker will be able to grow his private chain faster than the remaining network can grow the public chain.

Example 14. *This example is shown in Figure 3.5. Assume at the begin of the attack the longest chain contains blocks b_1 , b_2 , and b_3 . The attacker picks a recent block, e.g. b_2 and starts secretly extending b_2 . Once the attacker has extended his chain longer than the public chain he can publish it and all blocks in the public chain will be discarded.*

3.5.3 Selfish mining

In a selfish mining attack, the attacker does not violate the longest chain rule. Instead he violates the following mandate:

Publication mandate When a node finds a block, it should immediately announce this to the other processes.

The idea behind not publishing the newest block is that it denies the other nodes to try and extend the longest chain. Thus, other nodes waist resources, trying to extend a chain, that is not the longest chain.

For a detailed description and analysis of selfish mining see Chapter 26.1 of these Lecture notes form ETH Zurich.

Note 25. Results above show, that, if the attacker has more than $1/3$ of the hashing power, he can increase the ratio of blocks he creates in the blockchain by selfish mining.

Algorithm 3 Selfish mining

Idea: Mine secretly, without immediately publishing newly found blocks

Let l_p be length of the public chain

Let l_s be length of the secret chain

if a new block b_p is published, i.e. l_p has increased by 1 **then**

if $l_p > l_s$ **then**

 Start mining on b_p

else if $l_p = l_s$ **then**

 Publish secretly mined block b_s

 Mine on b_s and immediately publish new block

else if $l_p = l_s - 1$ **then**

 Push all secretly mined blocks

- If the attacker has more than $\alpha = 1/3$ of the hashing power, he can increase the ratio of blocks he creates in the blockchain by selfish mining.
- If the attacker has more than $\alpha = 1/4$ of the hashing power, and can reach $\gamma = 0.5$ half of the nodes before another miner can reach them, he can benefit from selfish mining.

3.6 P2P networking and network layer attacks

A node in bitcoin does not maintain connections to all 10.000 bitcoin nodes. Instead every node maintains a membership list, with addresses of other nodes. He maintains connections only to a few nodes selected at random from the list. We say that these connections form an overlay network.

In Bitcoin, nodes per default start by establishing 8 connections and extend this to up to 125 connections.

When a block is broadcast, every peer receiving the block first validates it and then forwards it to its neighbors.

3.6.1 Inventory messages and delivery denial attack

The forwarding of a block consumes significant bandwidth. Bitcoin therefore uses INVENTORY messages to announce the a block to neighbors. Receiving an INVENTORY message a node would request to receive the actual message from only one of its neighbors.

Nodes set a timeout when requesting a block. If they do not receive the block within the timeout, it is requested from a different source.

- In bitcoin version 0.10, the timeout for receiving a block is set to 20 minutes.

Definition 22. In a *delivery denial attack* a node does send INVENTORY messages, but when a block is requested, does not forward the block.

For extended details on this attack, see [Gervais et. al. in CSS'15].

Note 26. Due to the static timeout, a delivery denial attack can significantly slow down block propagation.

- Slowing delivery of blocks increases the probability of a fork, as can be seen from Theorem 2. This also increases the probability that a fork extends for several blocks.
- Slowing delivery of competing blocks may increase parameter γ in the selfish mining scheme and thus make selfish mining more appealing and profitable even when $\alpha < 1/4$.

3.6.2 Sybil attack

In the sybil attack, we assume that an attacker controls many IP addresses and machines, e.g. a botnet.

Definition 23. An attacker performing a *sybil attack* registers multiple peers to receive as many connections as possible. Additionally, he may spread false network addresses.

Note 27. A successful attack has the following **effect**:

- A successful attack allows the attacker to selectively reduce connectivity in the network.
- If the attacker chooses to cooperate in the forwarding of a specific message, block or transaction, this may spread significantly faster.
- Blocks or messages which the attacker chooses not to forward, may spread far slower.
- The increased network delay results in an increased fork probability.
- The selective connectivity gives the attacker an advantage, when performing selfish mining, i.e. increasing the γ parameter.

3.7 Attacks on transactions

It is possible to issue two transactions that both spend the same outputs. Only one of these transactions can be included in a chain, however, in case of a fork, each transaction may be included in a different branch of the fork.

Remember that a transaction counts as confirmed, if it is included in a block and a certain number of blocks is added on top of this block.

Definition 24. In a *double spend attack* on an *unconfirmed transaction* a payee accepts a transaction that is not confirmed, e.g. because he cannot wait for 20 minutes to sell a coffee.

The payer then issues a different transaction and tries to get this different transaction included in the chain.

Definition 25. In a *double spend attack* on a *confirmed transaction* a payee accepts a confirmed transaction.

The payer issues a different transaction and tries to get a fork including this transaction to become the longest chain.

Note 28. We now look at the different attacks from Section 3.5 and if they allow a double spend attack on confirmed transactions.

- A 51% attack allows to perform a successful double spend on a confirmed transaction. The attacker can simply create a secret fork including the second transaction, grow it to be longer than the public chain and publish once the original transaction is confirmed.
- During a selfish-mining attack it may happen that an attacker published a secret chain that is l blocks or longer. Especially, if the attacker already has created a secret chain that is 6 blocks longer than the public chain, he is sure to succeed in a double spend attack.
- A sybil attack may increase the network latency and thus cause forks. However, the probability for a long fork is still quite small.

Note 29. If attacker is doing selfish-mining, what is the probability that he currently is 6 blocks ahead (or more):

- If $\alpha = 25\%$ then about 0.001.
- If $\alpha = 33\%$ then about 0.01

These can be computed using the marginal probabilities from Chapter 26.1

3.7.1 Eclipse attack

The eclipse attack is a special case of the sybil attack. It is performed similarly, but rather than targeting the complete network, the target is a individual node or small group of nodes.

Definition 26. An attacker performing an *eclipse attack* registers multiple peers to receive as many connections as possible. Additionally, he may spread false network addresses to the victims.

The attacker aims to:

- I Monopolize all the connections of a single node or small part of the network.

Note 30. A successful attack has the following **effect**:

- A successful attack may allow the attacker to exclude individual nodes from the network. The excluded nodes may no longer receive new blocks and any blocks created by the excluded nodes will probably be discarded when the attack stops.

This allows the attacker to double spend a confirmed transaction. By creating a chain especially for the attacked node, which confirms the transaction. The main chain will be longer and contain the other double spend transaction.

3.8 Updating a blockchain

Most blockchain protocols are under constant change. There are bugs and security vulnerabilities that are fixed and new features or improvement proposals added.

As with other software, blockchain nodes do not accept and install updates instantly. For some controversial updates, nodes might even choose to not update.

Definition 27. An update that changes the rules, which blocks and transactions are valid, is called a *fork*.

Note 31. A software update that changes the code run by a node, but not its output is a *non-fork*. Here nodes implementing new and old versions seamlessly work together.

3.8.1 Soft fork

Definition 28. A *soft fork* is an update that makes some of the initial transactions or blocks, valid under a previous version invalid. However all blocks and transactions under the new version are valid according to the old protocol.

Example 15. A typical example of a soft fork is a security update, that rules out certain behaviors that were allowed on the previous version.

Note 32. The result of a soft fork depends on whether a majority of the nodes switches to the new fork.

- If the new version in a soft fork is accepted by less than 50% of the nodes, these will create their own chain.
- If the new version in a soft fork is accepted by more than 50% of the nodes, any block published on the old chain will eventually end in a short fork and be discarded.

3.8.2 Hard fork

Definition 29. A *hard fork* is an update that creates blocks that are not valid under the previous protocol. However blocks that are valid under the previous protocol are still valid under the new protocol.

This version of a hard fork is also sometimes called a *strictly extending hard fork*. It is generally considered easier to implement changes as a hard fork, than as a soft fork.

Example 16. An example for a hard fork is a new feature introduced, e.g. a new *Op code* for scripts. Blocks and transactions that do not utilize this new feature are still valid under the new protocol.

Note 33. The result of a hard fork depends on whether a majority of the nodes switches to the new fork.

- If the new version in a hard fork is accepted by less than 50% of the nodes, any block published on the new chain will eventually end in a short fork and be discarded.
- If the new version in a hard fork is accepted by more than 50% of the nodes, two chains will be created.

3.8.3 Hard and soft forks

Definition 30. A *hard and soft fork* is an update that creates blocks that are not valid under the previous protocol and also invalidates blocks from the previous protocol.

This version of a hard fork is also sometimes called a *bilateral hard fork*.

Note 34. In a hard and soft fork, there will always be two chains created.

3.8.4 Analysis

Forks that are caused by software or protocol updates have the potential to create forks of significant length. During such updates it may be easy to perform a double spending attack.

A safe variant is to require transactions to be included in both forks. However this is not possible for all variants, especially when spending outputs that were created in one of the forks.

There exist examples where a fork has resulted in two chains that were maintained separately.

Idea 5. A common idea is for nodes to first signal their readiness to switch to a new version, e.g. in the blocks they publish. This allows rules like:

- Only move to the new version if 95% of the last 2000 blocks have signaled to want to move to this version.

3.8.5 Pow as voting

The discussion in Section 3.8 shows that solving proof of work can be seen as majority voting, where every node has a voting share equivalent to the hashing power it is using to solve PoW.

It is possible to take this view also for forks that occur during normal operation. A node votes by trying to find a block extending a certain chain.

Note especially that this voting mechanism has a build in sybil resistance, since voting requires limited resources, e.g. CPU or AISCS cycles and electricity. It is not possible to create a higher voting share, by simply creating new identities on the system.

Chapter 4

Alternative PoW

4.1 Improving PoW

4.1.1 Alternative mining puzzles

See Chapter 8 of this book.

Definition 31. A suitable mining function has the following properties:

Adjustable difficulty It must be possible to adjust the difficulty to adapt to a growing network.

Fast verification Every node in the network needs to verify a solution. Thus verification should be easy, compared with computation.

Progress free The probability to solve the PoW function in the next second, should be independent of how long a process has been trying to solve it.

4.1.2 ASIC resistance

Asics resistance is motivated by the fact that bitcoin mining is currently done almost exclusively on specialized hardware, i.e. application specific integrated circuits (ASICs). There are currently produced by a single manufacturer, yielding a single point of failure and trust.

Further, these ASICs are difficult to use for other purposes than mining cryptocurrencies.

Definition 32. A PoW puzzle is *ASICs resistant* if specialized hardware can only provide a small benefit in computing this puzzle, compared to a general purpose CPU.

Note 35. Proposals for ASICs resistant puzzles include

Memory hard functions The idea for this is to design puzzles that require a lot of memory access. The idea is that memory access is harder to optimize than cpu computation, as done for hashing.

CPU benchmarks A recent proposal HashCore ICDCS'19 proposes to use CPU benchmarks, used normally during hardware development, to create PoW puzzles.

Both proposals result in functions that are harder to verify.

4.1.3 Proof of useful work

The idea behind proof of useful work, is that energy and hardware used to compute PoW solutions seems "wasted".

Definition 33. A PoW puzzle is *useful*, if its solution or computation has an application or utility outside of blockchain domain.

Note 36. Proposals for useful PoW are still subject to research. The main difficulty in using real world problems, i.e. from machine learning or protein folding, is to ensure progress freedom. Proposals include:

- Finding prime numbers
- Evaluate small degree polynomials: This can help to determine mathematical problems like vector orthogonality, shortest path problems, ...

Note that while finding problems, e.g. vector orthogonality have applications, it is usually not useful to simply find two random orthogonal vectors. Rather two orthogonal vectors should be found from a given set.

This poses the question, how the set, i.e. the problem should be submitted.

4.1.4 Proof of Storage

Definition 34. A *proof of storage* mining puzzle requires a node to store a certain file to be able to create a block.

Note 37. In a simple variant a proof of storage is a merkle inclusion proof. This can be combined with a classical PoW to adjust hardness, ...

Several problems arise:

- What file to store?
- How to distinguish a stored file from a file, retrieved if necessary.
- How to deal with the increased payload due to merkle proofs.

A probably better alternative is to use the amount of storage a node provides as stake, in a Proof of stake scheme.

4.2 Proof of Stake

Bitcoin uses PoW to avoid centralization. To perform a 51% attack, the attacker has to invest huge amounts of energy and hardware. The idea behind proof of stake is to use the cryptocurrency for this purpose. I.e. to do a 51% attack in PoS, the attacker needs to own 51% of the currency.

Definition 35. Proof of work distributes block rewards to miners, equivalent on the energy and hardware cost they have invested. Proof of stake aims to distributed mining rewards, depending on the amount of money (i.e. cryptocurrency) the miners have frozen for this purpose.

Example 17. In PPCoin (Peercoin) a miner identified by *addr*, that has deposited $\text{coin}(\text{addr})$ can supply the current block, if

$$H(\text{prevBlockHash}||\text{addr}||\text{timeinseconds}) < d_0 \cdot \text{coin}(\text{addr})$$

- Here d_0 is a base difficulty. The probability that a miner with a specific address *addr* can mine the next block is proportional to $\text{coin}(\text{addr})$.
- **timeinseconds** shows time in seconds. Thus a miner gets a chance to submit a solution every second.

This constructions gives several **limitations and problems**. Some of these are common for many PoS schemes.

Predictability A miner can predict whether he will be able to mine the next block. He can increase **timeinseconds** above the current time.

PoW next block A miner with sufficient resources can try to tweak the current block, s.t. he will be able to also mine the next block.

Non deciding In case of a fork, a miner does not have to decide on which block he wants to mine. It is feasible to mine on both chains. Thus, forks may prevail for long.

History rewrite Theoretically it is feasible to rewrite the complete, or a large part of the history of this chain.

In the following we present some mitigations for the above vulnerabilities. In Chapter 9 we will discuss more advanced techniques to address these issues.

Predictability To reduce the effect of predictability, it is important that nodes check the timestamp received on blocks created by others, and do not accept blocks published ahead of time.

PoW next block The ability for this attack significantly depends on what data is included in `prevBlockHash`. If this hash includes transactions, it is possible to change transactions, as in the attack.

A mitigation is to not include the transaction hash in `prevBlockHash` and instead only use the proof of the last block:

$$\pi_i = H(\pi_{i-1} || \text{addr} || \text{timeinseconds})$$

To still provably include the transactions in a block, it is then necessary to rely on signatures.

Non deciding The common technique to address this is slashing. Slashing is a technique that punishes misbehaving miners. The idea is, if some miner provably did work on two forks, then his depositum (stake) is discarded, or given to others.

History rewrite One technique for this is to include checkpoints, or similar and to not accept new branches, which diverge from the current branch before the checkpoint.

4.2.1 Fairness

Additionally to the above problems, research has shown that PoS, as in Peercoin is more likely to give unfair reward distributions than PoW. This subsection presents results from: Huang et al., "Do the rich get richer?", SIGMOD'21. Note that Huang et al. refer to our PoS (Peercoin) as multi-lottery (ML-PoS).

Definition 36. We can say that a reward system is fair, if every Miner/Participants receives rewards proportional to his invested energy or stake. However, reward distribution in a PoW or PoS system is probabilistic. Thus, we can consider fairness in two ways:

- **Expected fairness** expresses if the reward which a node is expected to receive, or receives over long time is fair.
- **Fairness variation** or robust fairness expresses the likelihood that the reward actually received deviates from the fair share.

Theorem 4. *According to Huang et al. both PoW and PoS, as in Peercoin have expected fairness, however, the fairness variation in PoS is significantly larger than in PoW.*

Chapter 5

Scaling the Blockchain

Bitcoin currently achieves about 7 transactions per second. Also, to be sure that a transaction is included, clients wait for 6 blocks, which takes on average 1 hour. In this chapter we discuss different approaches to increase this throughput and reduce confirmation latency.

Definition 37. **Throughput** of a blockchain is the maximum number of transaction that can be included in the chain per second.

Definition 38. **Confirmation latency** of a blockchain is the time it takes from when a transaction is first included in a block, until it can be considered confirmed, and a client can be sure that it will not be removed from the chain.

5.1 Blocksize and block interval

To increase this, it was heavily discussed, whether the blocksize should be increased, or block frequency can be reduced.

Definition 39. A change in maximum blocksize or target block interval is called a *reparametrization*.

Research, (see resources.md) has suggested that network latency grows linearly with blocksize.

Lemma 8. *A reparametrization that increases blocksize will result in higher network latency and thus in an increased fork probability.*

Proof. Theorem 2 says that increased network latency results in higher fork probability. \square

Lemma 9. *A reparametrization that decreases target block interval causes an increased fork probability.*

Proof. A reduced target block interval, i.e., less than the current 10 min value, will result in an increased probability that a block is found within a second (p in Theorem 2). \square

An increased fork probability creates the following **problems**:

1. **Security** Forks among the honest miners make selfish mining more efficient and make a 51% attack easier.
2. **Centralization** Forks are bad for small miners, since they are likely to lose their mining reward in a fork. Thus frequent forks favor centralization and the formation of larger mining pools.

To measure these, the following metrics exist: See [Bitcoin NG] for definition and evaluation.

Definition 40. Mining power utilization is the number of blocks in the longest chain, divided by the total number of blocks created.

Mining power utilization is related to the security problem above.

Definition 41. Fairness given the fraction with the largest mining power (or one such fraction). Fairness is defined as the percentage of block in the longest chain, not published by this party, divided by the percentage of blocks not mined by this party, relative to all blocks, including forks.

5.1.1 Ghost

GHOST is a proposal to avoid the security problem (1.) that may arise with more frequent forks, namely increased vulnerability to selfish mining and double spend attacks. This is especially relevant to maintain security after reparametrization for increased throughput.

Definition 42. The *Greedy heaviest-observed subtree (GHOST)* rule says, instead of selecting the longest chain the root of the subtree containing most blocks should be selected.

Note 38. • As long as only a single fork exists, the GHOST rule is identical with the longest chain rule.

- In a selfish-mining attack the attackers chain will not contain any forks, since only a single node/agent is trying to extend it. The chain build by honest nodes may show additional forks. Using the GHOST rule, these forks do not give an advantage to the attacker.
- Using network attacks the attacker may cause forks on the honest chain. Using the GHOST rule, these attacks have no impact on the probability of a successful attack.

5.1.2 Inclusive chains and uncles

Inclusive chains are a possibility to address Problem 2, **Centralization**. An increased fork probability may result in many blocks being discarded. This results in miners not getting their block rewards. It provides an incentive for miners/nodes to form larger groups, rather than mine individually. This is true also when using the GHOST rule.

Definition 43. The published blocks, both on the longest chain and in other forks, build a *tree*, routed at the genesis block. The *previous block hash* h_{-1} pointers point to the parent of a block.

Definition 44. In an *inclusive* blockchain, instead of just including the hash of the parent, a block b_{new} can include hashes of another block b , if:

- b is not an ancestor of b_{new}
- the depth of b (d_b) (i.e. distance from the genesis block) is less than the depth of b_{new} (d_{new}).

Blocks included as additional ancestors are called *uncles*.

Definition 45. In an inclusive blockchain, blocks may receive an **uncle reward** and **nephew reward**.

- A block b included as uncle in b_{new} receives a fraction of the block-reward, the **uncle reward**. The fraction reduces with the distance between b and b_{new} , i.e. $d_{new} - d_b$.
- A block b_{new} receives a small reward (**nephew reward**) for every uncle it includes.

Example 18. *Ethereum uses uncle blocks. An uncle b is rewarded $1 - \frac{(d_{new} - d_b)}{7}$ of the block reward. Thus uncles that are more than 6 blocks behind are not rewarded.*

For including an uncle a block creator is rewarded $\frac{1}{32}$ of a block reward.

Ethereum also requires that an uncle is a child to one of the ancestors of the nephew. Thus, on longer forks, only the first block can be an uncle.

Note 39. Security and incentives:

- The possibility to receive a block reward, being not included in the main chain, reduces the urge to form larger mining groups (pools).
- The reward an uncle receives plus the reward for including an uncle is less than the block reward. This should incentivize nodes to try and extend the main chain.
- The reward for uncles is reduced over distance. This discourages keeping blocks secret.

Theorem 5. *In an inclusive blockchain, the selfish mining attack becomes profitable at a lower α threshold than without inclusive mining.*

This can be leveraged to some extent, if the difficulty is adjusted not based on the frequency of blocks on the main chain, but based on the frequency of blocks on the main chain, and uncles.

Proof. See Selfish Mining in Ethereum




Figure 10. The profitable threshold of hash power in Bitcoin and Ethereum.

In the above figure from [Selfish Mining in Ethereum, ICDCS'19] Scenario 1 is where block difficulty is only adjusted based on the new blocks on the main chain. Scenario 2 adjusts difficulty based on creation of blocks and uncle blocks.

□

5.2 Uncles for scaling

In Ethereum, the transactions included in uncle blocks are not applied. However, it is possible to also apply these transactions.

Theorem 6. *It is possible to define a deterministic order on all blocks in a chain and all uncles in a chain, such that the order extends when the chain is extended. This allows to execute transactions that are included in the uncle blocks, but not in the main chain.*

Proof. The order is done based on the following principles, from highest to lowest priority:

- Ancestors of b_{new} are ordered before uncles of b_{new} .
- Uncles of b_{new} are ordered before b_{new} .
- Uncles are ordered according to their depth.

- At the same depth, uncles are ordered according to their hash.

□

Note 40. Inclusive blockchains have the following effect on scalability:

- If uncle blocks include transactions that are not included or conflicting with the main chain this can further increase scalability. Execution of transaction in uncles is not implemented in Ethereum.
- Given execution of transactions in uncles, it is difficult how to incentivize different blocks in a fork, to include different transactions.
- It is possible to extend inclusive blockchains to allow uncles that are not children of ancestors in the current main chain. Theorem 6 still holds, but becomes more complex.

5.3 Bitcoin NG

See resources.md on course info for Video, slides, paper, ...

Definition 46. Bitcoin-NG differentiates between *Keyblocks* and *Microblocks*.

Keyblocks Include a proof of work. No transactions, a public key, and the hash of the last Key- or Microblock.

Microblocks Include no PoW but transactions, a previous block hash, and a signature matching public key from last KeyBlock.

Fork resolutions Longest chain rule (or GHOST) is used but only looks at the Keyblocks. **Fee distribution** Fees are distributed between the creator of the microblock (40%) and the creator of the next Keyblock (60%).

Note 41. (Bitcoin-NG)

- Security and fork probability (keyblocks) is decoupled from throughput and transaction rate (microblocks).
- The distribution of fees (40/60) ensures that it is better to reference the latest microblock, rather than keeping the transactions for your own microblocks. Also significant incentive for actually publishing microblocks.

Problem

- Since Bitcoin-NG relies on a leader, progress can be reduced through leader failure. A single leader failure will not be a problem, but if an attacker causes multiple leaders to fail, e.g. by a DDOS (Distributed denial of service) attack, transaction throughput may be significantly reduced.

Example 19. To understand the choice of 40%, 60% fee distribution, consider the example in Figure 5.1. Squares represent key blocks, while circles represent microblocks.

In Figure 5.1 the minor of the red block gets 60% of the fees from the yellow transactions from microblock b'_0 .

If he instead mines on top of b_0 , he can himself publish a microblock including the yellow transactions. This is shown in Figure 5.2. In this case the red minor receives 40% of the fees from yellow transactions. Additionally, if the red minor has a fraction of the mining power equal to α , he has a probability α to mine key-block b_2 and get additional 60% of the yellow fees. If $\alpha < 0.33$ then $0.4 + \alpha \cdot 0.6 < 0.6$. Thus the red minor will try to extend b'_0 .

Similarly, if the blue minor does not publish block b'_0 but instead hopes to publish b_2 in Figure 5.2, his expected fraction of the yellow fees is $\beta \cdot 0.6 < 0.4$. Thus for a hashing fraction $\beta < 0.33$ it is better to publish b'_0 .




Figure 5.1: Red miner gets 60% of yellow transactions.




Figure 5.2: Red miner expects to get $40 + \alpha \cdot 60\%$ of the yellow transactions.

5.4 Sharding for blockchains

When sharding, the idea is to divide the blockchain into many smaller chains. Each small chain (i.e. shard) maintains part of the data.

Definition 47. In a *sharded blockchain* every node only maintains a part of the data and processes transactions that access the data he stores.

Note 42. **Problems for sharding:**

- A) How to distribute the state?

- B) How to process updates that access state in multiple shards?
- C) How to avoid that an attacker takes over one shard?

Solutions

- A) Using consistent hashing. E.g., shard 1 stores all outputs that belong to addresses that start with 01.
- B) There are different cases:
 - Payments from one shard to another, can be split into a withdrawal and deposit. Execute withdrawal first. To execute deposit, must reference withdrawal.
 - Transactions that are conditioned on multiple shards use a variant of 2 Phase Commit, e.g. first lock required outputs. When all necessary outputs are locked, execute transaction.

Note that variant 2 can diminish the benefit of sharding, since it requires coordination and locks state.

- C) There are two ideas:
 - Distribute miners to shards according to their hashes, and redistributed frequently.
 - Allow miners to participate in multiple shards, and to publish blocks on multiple shards with a single proof of work.

Both of these are problematic. In the first, redistribution is reducing the benefit of sharding. The second encourages miners to form pools.

Chapter 6

System models

In the Bitcoin protocol, any node can join. Additionally any entity can create multiple accounts on bitcoin. Accounts are per default anonymous. This model is called **unpermissioned**. No permission is needed to participate.

6.1 Unpermissioned systems

In an unpermissioned or permissionless system anyone can join, usually anonymity is provided. Thus the creation of sybils is usually possible.

To mitigate the existence of sybils, unpermissioned blockchains use Proof of Work or Proof of Stake or similar mechanisms.

6.2 Permissioned systems

Permissioned systems assume there exists a list of members. New members need permission or authentication via a central authority to join. Table 6.1 gives an example of a table of members.. Typically assume that all members have a copy of the membership list.

The existence of a membership list, or possibility to create and distribute one is often referred to as public key infrastructure (PKI). It allows any two members to contact each other (using listed IP addresses), and create an authenticated channel (using public keys).

ID	<i>pubkey</i>	IP	...
1	0x1f3		
2	21xf3		

Table 6.1: Membership list example.

In a permissioned system the central authority prevents sybils.

- Example 20.** 1. *The classical example for a permissioned blockchain is a blockchain run together by a set of organizations. Typically, we assume that each organization runs a node. The Libra/Diem blockchain is an example for this. For Libra, a non-profit central organization exists, which is functioning as central authority.*
2. *If a single system, run within one organization is running on multiple servers, this also is a permissioned system. The different servers have unique ids and the system administrator functions as central authority.*
3. *A more open example is a system that requires nodes to authenticate with Bank-ID, passport or sharing their social media activity. Here the central authority is the issuer of Bank-ID, passport or social media.*

6.3 Failure models

We distinguish different failure models, based on the failure assumptions, e.g. what failures may happen. All apply to permissioned systems, some are applicable to unpermissioned systems.

No failure Assuming that members do not fail, it is possible to use authenticated channels for trusted interactions.

Central node does not fail Given a centralized node, that processes trust to not fail, they can build a centralized system, where the trusted node is responsible to manage state and interaction for the other members.

This model is used in cloud services, where all users connect to the service provider and trust this provider to store and maintain their application state and coordinate interactions with other users.

Crash failure In this model, nodes can stop by crash failures. It is not possible to communicate with a crashed node and the data stored on a crashed node may be lost.

This model is often assumed for machines that cooperatively run a service within a trusted domain, e.g. servers within a data center.

In this model it is possible to implement consistent services, e.g. a blockchain, that work as long as a majority of the nodes is running. Examples are the Paxos algorithm taught in DAT520.

In this model, as shown by the Paxos algorithm it is possible to develop systems that simply halt, when the failure threshold is violated. Thus, even if a majority of nodes fail, nothing bad will happen.

Byzantine fault tolerance BFT It is assumed that *any node may fail or misbehave*, i.e. a faulty node may not only stop, but may act maliciously, trying to sabotage the application. However, it is assumed that *only a small fraction of the nodes actually fail or misbehave*.

Misbehavior may for example be caused by

- virus or malware
- misconfiguration
- sabotage

The above assumption says, that at most a small fraction of the nodes will be victim to any of these.

In this model it is possible to implement consistent services, e.g. a blockchain, that works as long as a large majority, e.g. $2/3$ of the members are correct.

Selfish or rational misbehavior All nodes could misbehave if it benefits them. Such nodes are called *selfish* or *rational*. Different from the BFT model, we do not assume a failure threshold, but we assume, nodes will not misbehave if this is not in their interest.

In this model it is possible to design algorithms using game theory. For this, every node is assigned a utility function:

For example the utility of participating in bitcoin mining is the block reward and transaction fees, minus costs for computation and networking, e.g. energy and hardware.

The goal for game theory is to show that nodes cannot increase their utility, by deviating from a protocol, e.g. by deviating from the longest chain rule.

Definition 48. In game theory we assume that agents or players are free to choose from different actions. The set of actions an agent takes are called a **strategy**.

In traditional game theory we assume that games are deterministic, meaning that given a set of strategies, one for every agent, we can determine what will be the utility of every agent.

Definition 49. A **Nash-Equilibrium** is a set of strategies, one for every agent, such that no individual agent can improve his utility, by choosing a different strategy.

For a protocol it is desirable that for every node to follow the protocol is a Nash-Equilibrium. This means that agents will not profit in utility by deviating from the protocol.

However, note that the notion of Nash-Equilibrium does not exclude the possibility that coalitions may successfully deviate.

Note 43. There exists variants of the Nash-Equilibrium that take into account, e.g. the possibility for small coalitions, the fact that agents may only deviate if it gives a significant benefit, or the fact that a game may be probabilistic.

Chapter 7

A BFT blockchain protocol

7.1 Proof of certification

Unpermissioned blockchains require that published blocks carry a nonce that causes the block hash to meet a specific difficulty (e.g. start with a specific number of zeros). This proof of work has two main functionalities:

1. **Publishing rate:** Requiring a proof of work ensures that blocks are published at a limited rate. This ensures that, whp., a block is propagated throughout the network, before the next block is published.
2. **Fork probability:** If, while a block is propagated throughout the network, another block is published, these blocks create a fork and put the system in an undecided state. Proof of work ensures that the probability that two blocks are found concurrently is small.
3. **Detectable system split:** If a subset of the network and miners of a PoW blockchain splits off, to form their own chain, you can detect that by the drop in block rate, since the smaller network will not be able to solve PoW at the same rate.

In a permissioned system, similar guarantees can be achieved using a certificate.

Definition 50 (System Model). A permissioned system is comprised of N nodes n_1, n_2, \dots, n_N . We assume that nodes have access to digital signatures, and that each node knows public keys of all other nodes.

Definition 51. A certificate for a block b is a collection of digital signatures for b . A certificate contains signatures from more than $N \cdot \frac{2}{3}$ different nodes. We write c_{min} for the minimum number of signatures contained in a certificate

$$c_{min} = \left\lceil \frac{2N + 1}{3} \right\rceil$$




Figure 7.1: Blocks in a tree and depth d of blocks.

Idea 6. We now give an intuition how certificates can limit publishing rate and fork probability:

1. **Publishing rate:** To publish a block with certificate, this block has to be transmitted, validated and signed by at least c_{min} nodes. Thus, blocks cannot be published faster, than (most) of the nodes can validate them.
2. **Fork probability:** If less than $\frac{N}{3}$ nodes signed multiple conflicting blocks, no two conflicting blocks can both receive a certificate.
3. **Prevent system split:** If a subset of the nodes are split of from the main network, they will not be able to create any certificates.

7.2 Safety

In the following we define a set of rules that correct nodes should follow. We assume that at least c_{min} nodes are following these rules.

The *previousBlock* pointers create a tree structure on blocks. We can thus refer to the *depth* of a block. We write $b.depth$ for the depth of block b . Figure 7.1 shows a tree with different depth levels. Here $b'_4.depth = 4$. Similarly, we refer to the parent, ancestors or descendants of a block. In Figure 7.1, e.g., b''_5 is a descendant of b''_3 and all blocks are descendants of b_0 . Further, b'_3 is the parent of b'_4 and b_2 is an ancestor of b'_4 .

We can not define the first rule.

Rule 1. After signing a block at depth d , only sign at depth $d' > d$.

The first rule says that nodes should only sign at increasing depth. Especially, they should not sign at the same depth twice.

As noted in Section 7.1, to prevent forks we should also add a rule that prohibits changing from one branch of the tree to another, e.g., signing b''_5 in Figure 7.1 after signing b'_3 and b'_4 . We further note, that it is possible, that




Figure 7.2: Blockes with red justification link, confirm block b .

no block at a certain depth receives a certificate. E.g., b'_3 and b''_3 may both be signed by 5 out of 10 nodes. Thus, no block has the required ($c_{min}(10) = 7$) signatures. Therefore it should still be allowed for a node to sign b'_4 after signing b''_3 .

Definition 52. The *locked block* at node n_i , $n_i.lock$ is the block b at highest depth, such that n_i has (or has seen) a certificate for b .

Rule 2. A node n_i only signs a block that is a descendant of $n_i.lock$.

We note that if c_{min} nodes sign a block b , that does not imply that these nodes know the certificate of this block. Certificates have to be collected and disseminated by some node. We therefore define the following:

Definition 53. We assume that every block b_i contains a certificate for a block b_j such that b_j is an ancestor of b_i . We say that $b_i.justify = b_j$.

Lemma 10. For some node n_i that follows Rule 2, it holds that, after signing block b , $n_i.lock.depth \geq b.justify.depth$ holds.

Example 21. In Figure 7.1, if $b'_4.justify = b'_3$, then after signing b'_4 a node following Rule 2 will not sign b''_5 .

Definition 54. We say that a block b is *confirmed*, if there exist blocks b' and b'' , such that, $b = b'.justify$, $b' = b''.justify$, and $b.depth = b''.depth - 2$.

We note that the notion of a confirmed block is similar to proof of work blockchains like bitcoin, where a block counts as confirmed if it has been extended by a certain number of blocks, e.g. 6 blocks in bitcoin.

Theorem 7. If nodes follow Rule 2 and a block b is confirmed, then any certified block at depth $d > b.depth$ is a descendant of b .

Proof. If b is confirmed, there exist b' and b'' as in Definition 54. b'' contains a certificate for b' . Thus at least c_{min} nodes have signed b' . b' contains a certificate for b . Thus upon signing b' , a node n_i sets $n_i.lock = b$.

We have to show that every certified block β with $\beta.depth > b.depth$ is a descendant of b . We proof this by induction over $\Delta_d = \beta.depth - b.depth$.

If $\Delta_d = 1$ then $\beta = b'$ since at most one block at depth $b.depth + 1$ can be certified.

Let $\Delta_d = n + 1$. Both β and b' are signed by c_{min} nodes. Let n_i be a correct nodes (following Rules 1 and 2) that signed both β and b' . Due to Rule 1, n_i signed b' before β . Thus $n_i.lock = b$ did hold. The induction hypothesis implies that when signing β , $n_i.lock$ was either b or a descendant of b . Rule 2 implies that β is a descendant of $n_i.lock$. Thus β is a descendant of b . \square

7.3 Liveness

In Section 7.2 we have defined rules, how correct nodes should sign blocks and how we can identify confirmed blocks, based on published blocks.

However, we did not define when and how blocks should be created. Further, we note that if at every depth, many blocks are created, it is possible that no block will ever receive a certificate.

To resolve this, we assign leaders to every depth and only allow the leader to publish blocks at his depth:

$$\text{leader}(\text{depth}) = n_{\text{depth} \bmod N}$$

Algorithm 4 shows how nodes only sign blocks at the next depth and only if the block is published by $\text{leader}(\text{depth})$. However, based on a timeout, nodes can hop over one depth.

Algorithm 4 Rotating leader

```

1:  $\text{depth} = 1$ 
2:  $\text{timer} = \text{start}()$ 
3: on receive  $b$  from leader( $\text{depth}$ )
4:   if  $b.\text{justify}.\text{depth} > \text{lock}.\text{depth}$  then
5:      $\text{lock} = b.\text{justify}$ 
6:   if  $b.\text{depth} = \text{depth}$  and  $b$  descendant of  $\text{lock}$  then
7:      $\text{sign}(b)$ 
8:      $\text{depth}++$ 
9:      $\text{timer} = \text{restart}()$ 
10: on  $\text{timer}$  finish
11:    $\text{depth}++$ 
12:    $\text{timer} = \text{restart}()$ 
13: on leader( $\text{depth}$ ) =  $\text{self}$ 
14:   ask all nodes for locked certificate
15:   if block is missing at depth  $\text{depth} - 1$  then
16:     create empty block at  $\text{depth} - 1$ 
17:   create block at  $\text{depth}$  including deepest certificate

```

Lines 13 to 17 show how a correct node should propose a block. First

a correct node needs to query other nodes, especially his predecessor, for certificates they have collected. There are two cases:

- a) If a nodes collects a certificate for the last block ($depth - 1$), he can immediately publish a new block including this certificate.
- b) If a nodes does only collect certificates for older blocks. He includes the certificate at highest depth in his block. If no block at $depth - 1$ is known to the node, he may also create that block.

The situation from Case b) is also shown in Figure 7.3.




Figure 7.3: Leader for $b_f.depth$ may also create b_e .




Figure 7.4: Figure illustrating Example 22.

Example 22. Assume $N = 4$. Thus $c_{min} = 3$ and we have nodes n_1 , n_2 , n_3 , and n_4 .

In Figure 7.5, green subscript shows the nodes that have signed a specific block. Nodes n_2 , n_3 , and n_4 have signed b'_3 , while n_1 has signed b''_3 . Further, n_2 has signed b'_4 . Thus $n_2.lock = b'_3$. We assume that only n_2 has seen the certificate for b'_3 .

Since n_3 and n_4 have not seen a certificate for b'_3 they have, together with n_1 signed b''_4 . n_1 has seen this certificate, when signing b''_5 . Thus $n_1.lock = b''_4$. We note that in this example, none of the nodes have behaved faulty.

Example 23. This example extends Example 22. Thus we again assume $N = 4$. In Figure 7.5 shows a situation as in Example 22.

Node n_2 has locked b'_3 and node n_1 has locked b''_4 . We note that Rule 2 prohibits n_2 from signing b''_6 and n_1 from signing b'_6 . This may deadlock the system.

However we note that Rule 2 allows n_2 to sign b_5'' as shown in Figure ?? . The reason is that b_5'' includes a certificate for b_4'' . This allows $n_2.lock$ to be updated.

In Example 23, it is not helpful, to let n_1 , n_3 and n_4 report whether they have voted for b_4' or b_4'' . The reason for this is that one node may violate Rule 1. Thus n_3 or n_4 may also sign b_4'' .




Figure 7.5: Figure illustrating Example 23.

Existing systems show two possible solutions for the problem shown in Example 23. The PBFT algorithm allows n_2 to sign b_6'' if he receives messages from n_1 , n_3 and n_4 saying that b_2 is their last certificate.

Other systems, like Tendermint, include a long timeout while collecting certificates on Line 14, during which the new leader waits for additional certificates. This ensures in Example 23, if n_1 is correct, the certificate for b_4'' will be forwarded to the new leader.

7.4 Safety revisited

In the following we present a variant of Rule 2 that does not lead to the problem from Example 23 in Section 7.3.

We first define a 3-locked block.

Definition 55. A node n_i sets $n_i.lock_3 = b$, if b is the block at highest depth, s.t. n_i has signed a block b'' with $b''.justify = b'$ and $b'.justify = b$.

Rule 3 (replaces Rule 2). A node n_i signs a block b only if

- a) b is a descendant from $n_i.lock_3$
- b) $b.justify.depth > n_i.lock_3.depth$.




Figure 7.6: Block b_a is locked on signing b_f .




Figure 7.7: Showing 3-locks that allow two sign any block extending b_2 .




Figure 7.8: Rule 3 b) allows n_2 to sign b''_7 .

Example 24. Figure 7.7 shows position of 3-locks in the same situation as shown in Figure ???. Rule 3 allows to sign any descendant of b_2 .

Figure 7.8 shows a similar situation as in Figure 7.5, only using 3-locks. n_2 has signed b'_6 and set its 3-lock to b'_3 . The leader for depth 7 failed to collect the certificate for block b'_5 from n_2 . Thus he extended block b'_4 . Rule 2 b) still allows n_2 to sign b''_7 .

Using Rule 3, it is enough if the leader waits for the last certificate from c_{min} nodes on Line 14 of Algorithm 4.

Definition 56. We say that a block b is *3-confirmed*, if there exist blocks b' , b'' , and \hat{b} such that, $b = b'.justify$, $b' = b''.justify$, and $b'' = \hat{b}.justify$ and $b.depth = b''.depth - 2$.

Theorem 8. If nodes follow Rule 3 (instead of Rule 2) and a block b is




Figure 7.9: Block b is 3-confirmed.

3-confirmed, then any certified block at depth $d > b.\text{depth}$ is a descendant of b .

Theorem 9 (Speculation). *Using Rule 3, if a node n_i cannot sign a block b due to Rule 3, then n_i has a proof of misbehavior against the leader that published b .*

Chapter 8

Smart Contracts

For additional information on Ethereum, especially on the organization of storage and use of trees, see *Tabatabaei et. al.*, <https://doi.org/10.48550/arxiv.2207.02264>.

8.1 Ethereum and Smart Contracts

Ethereum is a Proof of Work blockchain, very similar to BitCoin. Ethereum has its own cryptocurrency, *ether*, and uses hashes of public keys as addresses and identities, as bitcoin does. However Ethereum implements several of the optimizations mentioned earlier.

- Ethereum has an average block delay of only 12 seconds.
- Ethereum uses a different P2P network, namely Kademlia.
- Ethereum uses a different PoW function, Ethash, which aims for ASIC resistance, as discussed in Section 4.1.2.
- Ethereum uses Uncles and the GHOST rule, as discussed in Chapter 5.

Note 44. In 2022, Ethereum changed from PoW to a Committee based PoS. This did not involve significant changes for the SmartContract infrastructure explained below.

8.1.1 SmartContract

Ethereum is mainly known for enabling Smart Contracts. These allow to encode more complex logic than the scripts in Bitcoin, i.e. redeem conditions.

Definition 57. A **Smart Contract** is similar to an object in OOP. It has fields (state) and methods (functions) and is stored on the blockchain. Further, the contract is characterized by the following properties:

1. The code (function implementation) is public and cannot be changed.
2. Any user can invoke the functions and thus update the state.
3. The contract object can own money (ether)
4. Users can send money to the contract together with invoking functions.
5. The contract can send money to other contracts or users, inside functions.

Note 45. Points 1 and 2 above ensure that strangers can use the same contract. If they verify the contract code, they don't need to trust the user that published the contract.

Points 3- 5 allow to couple sending of money with function invocation. This simplifies the integration of money flows into an application.

8.1.2 Accounts

Ethereum does not use the UTXO model. Instead, Ethereum uses accounts. Ethereum distinguishes between external and internal accounts. For every account, the Ethereum nodes store the fields, **address**, **nonce**, **balance**, **storage root**, and **code hash**.

Definition 58. In Ethereum **external accounts** are accounts belonging to users or rather identified by a public key. For every external account, the following data is stored:

address The accounts identified, a hash of the public key.

balance The amount of ether owned by this account.

nonce The sequence number of the last transaction sent from this account.

Also external accounts have the fields **storage root** and **code hash**, but they are empty.

Definition 59. A **internal account** or **contract account** is the address of a smart contract. It has the following fields:

address The accounts identified, derived from a hash of the contract creator address and the nonce of the transaction that created the contract.

balance The amount of ether owned by this account.

nonce The number of other contracts (accounts) created from this account.

storage root The root hash of the storage, or state of the contract.

code hash The hash of the contracts code.

Note 46. Internal accounts have a storage root and code hash. However, while transactions can update the storage root, the code hash cannot be updated.

There are proposals to include the code hash in the derivation of the address. However, the code hash alone is not sufficient for derivation of the address, since two contracts may have the same code.

8.1.3 Transactions

Transactions in Ethereum are always invoked from an external account. Transactions can transfer money, invoke functions on smart contracts or both.

Transactions have the following fields:

nonce next sequence number for sender account

gas price how much ether the sender is willing to pay per gas unit

max gas the maximum gas consumption of the transaction

recipient destination Ethereum address

value amount of ether sent along with the transaction

data payload send with transaction, e.g. function identifier and arguments

signature the signature from the sender, including his public key

If the recipient is an externally owned account, the transaction simply transfers the *value*. If the recipient is a contract account, the *data* is decoded to detect which function should be run, and with what parameters. If no function is matched, the default function is run.

As part of the function execution, a contract can call functions a different contract. This is called an *internal transaction*. Internal transactions are always part of an external transaction.

Transaction validation When validating a transaction the following checks are done:

- The *nonce* is the next sequence number for the sender.
- The sender has sufficient funds to pay *value* and *fees*.
- The signature is correct.

Note 47. When validating a transaction, it is not checked that the transaction executes correctly. A function invocation that causes an error may still be part of a valid transaction, be executed and will cost fees.

In a smart contract, you can access the address of the sender as `msg.sender`. This allows the contract to authenticate the sender, relying on the signature check on the transaction.

Note 48. On internal transactions `msg.sender` will not indicate the address of the contract invoking the internal transaction.

`msg.origin` is available to show the sender of a transaction, to which an internal transaction belongs. However using this field may have security issues.

8.1.4 Gas

Fees in ethereum are calculated in *Gas*. The sending of a transaction, function invocation and every assembly operation has a specific (fixed) gas cost. When issuing a transaction, the sender specifies *Gas price*, saying how much ether he is willing to pay per Gas unit.

Thus executing the same transaction costs the same amount of gas, but the actual fee may differ depending on the gas price.

Similar to Bitcoin, miners will include the transactions with the largest gas price, since also block size is given in gas.

Note 49. Requiring gas prevents users from staging a DOS attack by writing an infinite loop.

gas limit A transaction specifies a *gas limit*, saying how much gas the transaction may use at most. If the gas limit is reached, the transaction is reverted. However, the transaction will still be included in a block and executed, and thus will cost the sender fees.

Note 50. The *gas limit* can be used to avoid unexpectedly high fees. The reason this is necessary is that other transactions may change the state of a contract, and state changes may affect the execution of a function/transaction.

Example 25. *Assume a function iterates over a list. If other transactions insert elements into that list, executing the iterating function will be more expensive.*

gas tuning and storage Since the invocation of transaction costs gas, it is common to optimize the gas usage of functions, since this will save users actual money.

Operations that need to allocate storage are extra expensive. That is because the values have to be stored by all Ethereum nodes, possibly forever. Similar, if some elements are removed from storage, this actually costs a negative amount of gas.

8.2 Smart Contract Security

See slides and <https://github.com/ethereumbook/ethereumbook/blob/develop/09smart-contracts-security.asciidoc>.

8.3 Oracles

Smart contracts do not have access to data hosted outside of the blockchain. This is significantly limiting their usefulness, since many application have the need to access real data. For example

- a sports-betting application needs access to match results.
- a financial application needs access to exchange rates.
- an insurance application needs to know wether an event occurred.
- an online-store application needs access to wether a good was shipped.

For some kind of information, e.g. exchange rates, weather data, ... trusted sources and web-apis exist. However it is not possible for a smart contract to make a https request.

This problem can be solved by using an oracle:

Definition 60. An **oracle** is a special smart contract that published real world data onto the blockchain.

Example 26. Assume we want to build an application, that allows event hosts (e.g. festivals) to ensure against bad weather. Our application would like to access both historical weather data and predictions from a public api like *yr.no*.

Figure 8.1 shows the different components of the application. It has the following components:

- The insurance contract includes application logic of our insurance.
- The oracle contract can be called from the insurance contract, when access to the web-api is needed. The oracle contract will then emit an event.
- A web application, oracle web-app listens for events from the oracle contract. Upon events this application invokes requests to the public api. It then invokes a function on the oracle contract with the requested data. The oracle contract calls a callback on the insurance contract.




Figure 8.1: Oracle example.

Note 51. An oracle is in itself a trusted source. In the architecture described in Example 26, the insurer and its client both need to trust the weather service, but also the oracle web app. That is because the oracle web-app may falsify the information received from the weather service.

The oracle contract is separated from the insurance contract, so that it can be replaced, if the oracle seems not trustful.

Improving trust There exist measures that can be used to increase trust in an oracle.

1. Vote among multiple oracles.
2. Use trusted hardware to run the oracle web-app, i.e. Intel SGX.
3. Use zero-knowledge proofs (checked by oracle contract) to show that information was correctly received using https.

Point 1. has the advantage that it can also be applied, if no trusted web-api is available. Points 2 and 3 are interesting, since they also allow to proof facts about private data, without revealing that data. In that case, the oracle web-app can be run by the owner of private data.

8.4 Off chain transactions

The idea behind off-chain transactions is to allow two parties to interact outside of the blockchain, i.e. without paying fees or waiting for long confirmation times. Only if the parties disagree on transactions, they use the blockchain to resolve the dispute.

This technique can mitigate the following problems:

Fees blockchain transactions incur high fees. They can therefore not be used for small, every day, transactions.

Scalability the blockchain has a limited throughput and can therefore not handle all our every day transactions.

Confirmation time blockchain transactions have long confirmation times, e.g. 1hour in bitcoin. They are thus not useful for everyday transactions.

8.4.1 Uni-directional payment channel

A unidirectional payment channel allows a payer A to repeatedly transfer money to a payee B off chain.

- To open the channel, a smart contract is deployed. A funds the contract with an initial balance b .
- A can then send money to B by simply sending B a signed statement, including a balance $b_B \leq b$.
- A can repeatedly send money to B by sending another signed statement with an increased balance $b'_B > b_B$. However $b'_B \leq b$ must hold.
- B can checkout his money, submitting the last balance b'_B , that he received from A to the contract.
- Receiving the correctly signed balance b''_B from B the contract will payout b''_B to B and $b - b''_B$ to A .

The contract also has an expiration date. If B does not submit a signed balance, before expiration, then A can submit a transaction, terminating the contract and refunding b to A .

Note 52. The benefit of an uni-directional payment channel is that it incurs fees only when opening and closing. Also, once the channel is opened, transfers can be confirmed by B instantly, without waiting for the blockchain.

The drawback of uni-directional payment channels are:

- The channel is limited by its initial balance.
- The funds from A are locked inside the channel.

8.4.2 Bi-directional payment channel

A bi-directional payment channel allows multiple off chain payments between A and B , sending funds back and forth:

- To open the channel, a smart contract is deployed. A funds the contract with balance b_A and B funds the contract with balance b_B .
- To send c funds to B , A creates a signed message with balances $b_A^1 = b_A - c$, $b_B^1 = b_B + c$ and an increasing sequence number $nonce = 1$.

$$\langle (b_A^1, b_B^1), 1 \rangle_{\sigma_A}$$

- A and B can repeatedly send money, by changing balances and increasing the nonce.
- To close the channel, B can submit a message signed by both A and B

$$\langle (b_A^i, b_B^i), i \rangle_{\sigma_A, \sigma_B}$$

When this happens, a predefined timeout is set.

- Before the timeout expires, A can submit a balance with a higher nonce.
- After the timeout expires, A and B can withdraw their balance b_A^i .

Note 53. The bi-directional payment channel allows to send money back and forth. If the amount transferred from A to B is similar to the amounts transferred from B to A , then many transactions can be sent back and forth.

One drawback of bi-directional channels is the need to define a timeout. A small timeout may be difficult to meet. Especially since blockchain transactions may see a lot of contention. If timeouts are long, funds are locked. Additionally both parties need infrastructure, to survey the blockchain and react within the timeout if necessary.

8.4.3 Payment networks

While payment channels can allow cheap transfer of funds, channels lock a users funds and the creation of channels is still subject to fees and the long confirmation time of the blockchain.

Therefore, creating channels with all users, that we send money to is not feasible.

Payment networks allow to transfer funds across multiple channels. Thus A can send funds to B without creating a channel, if the graph created by all channels includes a path from A to B .

Example 27. Assume user A maintains a channel to X , and user B maintains a channel to Y . X and Y could for example be larger facilitators, e.g. banks.

Further we assume that a channel exists between X and Y . A payment can now be forwarded from A , via X and Y to B .

To transfer fund from A to B one can ask Y and X to forward such a payment. The smart contracts can be designed, such that, if X promised to forward a payment, and received a payment from A , then forwarding to Y can be enforced.

Chapter 9

Hybrid blockchains

Blockchains like bitcoin, running in an unpermissioned setting only provide probabilistic guarantees. This applies both to proof of work, proof of stake or proof of utility based systems. This is different from BFT systems used in permissioned settings, that provide strong guarantees.

9.1 Consensus guarantees

In this section we investigate the different guarantees given and techniques to build a hybrid system, providing strong guarantees in an unpermissioned setting.

9.1.1 Bitcoin guarantees

According to Bitcoin-NG, the consensus mechanism deployed by Bitcoin (i.e. Nakamoto Consensus) gives the following guarantees assuming that Byzantine or misbehaving nodes hold at most $f < \frac{1}{4}$ of the mining power:

Termination There exists a time difference function $\Delta(\cdot)$ such that, given a time t and a value $0 < \epsilon < 1$, the probability is smaller than ϵ , that at times $t', t'' > t + \Delta(\epsilon)$ a node returns two different states for the machine at time t .

Agreement There exists a time difference function $\Delta(\cdot)$ such that, given a value $0 < \epsilon < 1$, the probability that at time t , two nodes return different states for $t - \Delta(\epsilon)$ is smaller than ϵ .

Validity If the fraction of mining power of Byzantine nodes is bounded by f , then the average fraction of state machine transitions that are not inputs of honest nodes is smaller than f .

Note especially that Termination and Agreement only hold probabilistically. This may not be a problem in practice, i.e., ϵ can be chosen small

enough that the difference between probabilistic and absolute guarantees can be ignored. However, to achieve a small ϵ may require long waiting times, as in bitcoin where the confirmation of a block may take more than 1h. Similarly, if one of these properties, e.g. Termination, is violated it is difficult to say whether this is due to the probabilistic nature of this property, or to the fact that assumptions on the failure threshold f were violated.

9.1.2 BFT guarantees

In comparison, we now analyze agreement and termination properties of the BFT protocol described in Chapter 7. Here we assume that at most $f < \frac{1}{3}$ of the nodes in a permissioned membership misbehave (are byzantine).

We have shown, that once a block is confirmed according to Definition 54 or 56 then all future certified blocks will be descendants of this block. We note that, if we consider a block as confirmed according to these definitions, we can also consider all its ancestors as confirmed. Especially, we get the following agreement and termination property.

Termination safety If a correct node considers a block at *depth* l as confirmed, it will never change this block.

Agreement No two correct nodes will disagree on which block is confirmed at *depth* l .

Termination liveness For some *depth* l , some block at depth $l' \geq l$ will eventually get confirmed.

For Termination liveness to hold, it is necessary that a sequence of blocks is proposed by correct nodes.

9.2 Hybrid BFT

Several blockchain systems propose to use concepts from permissioned BFT to improve Termination and Agreement properties of unpermissioned blockchains.

There are two goals:

- **Non-probabilistic termination:** The question is how to get a non-probabilistic termination guarantee, like termination safety in BFT instead of probabilistic termination guarantee.
- **Confirmation time:** Probably the main goal for hybrid blockchains is to reduce confirmation time. In Bitcoin, people wait for about 1h for a transaction to be confirmed. On the other hand, in HotStuff, the three blocks needed to confirm a transaction can be broadcast and signed within seconds.

Resources

ByzCoin is a research system from EPFL that builds BFT on top of a proof of work blockchain.

Casper FFG is a proposal to add a BFT based system, also utilizing Proof of Stake, on top of Proof of Work in Ethereum, to ensure non-probabilistic termination.

9.2.1 ByzCoin

The difficulty in running a BFT protocol, as proposed in Chapter 7 in an unpermissioned setting is to identify the set of nodes and avoid sybil attacks. Proof of Work or Proof of Stake can be used for this purpose.

In ByzCoin builds on Bitcoin-NG, but instead of just the leader signing microblocks, ByzCoin requires microblocks to be confirmed using a BFT protocol, i.e. collects a certificate for microblocks.

In ByzCoin the nodes running the BFT protocol and signing the microblocks are also determined by key-blocks. ByzCoin uses a fixed window size. The finders of the last Δ key-blocks participate in the BFT protocol. We note that in ByzCoin, nodes are not equal. Instead every node has a *voting power* relative to the number of PoW puzzles (key-blocks) he submitted in the current window. Thus, instead requiring a certificate with signatures from $\frac{2}{3}$ of the nodes, ByzCoin requires signatures from nodes, that together hold $\frac{2}{3}$ of the voting Power in the current window.

Note 54. • Using a sliding window ensures that nodes that become inactive are removed from the BFT protocol.

- Forks present a problem, since they give disagreement who is participating. This can be solved by requiring that a key-block is confirmed by δ (e.g. $\delta = 6$) other key-blocks, before the issue is included in the window.

Example 28. *Figure 9.1 shows an example of the window and voting power in ByzCoin.*

Note that when voting for block i nodes a and c together hold $2/3$ of the voting power. This is no longer the case in window $_i$.




Figure 9.1: ByzCoin: An example for the Byzcoin sliding window. Nodes that have mined a block within the window get to sign the certificate for the new block (i) mined by y . Including i we move to window_i

node	voting power in window	voting power in window_i
a	3	2
b	1	1
c	2	2
x	1	1
y	0	1

Table 9.1: Table showing voting power for Figure 9.1.

9.2.2 Casper FFG

Casper FFG, requires nodes participating in the BFT protocol to deposit "Stake", i.e. Ether. A node depositing above a minimum of Ether may participate in the BFT protocol. Again, a node has a *voting power* representing the fraction of the total "Stake". Casper FFG uses Proof of Work to avoid the leader election problem. I.e. proposed blocks need to contain a proof of work. This limits the number of blocks that can be proposed concurrently.

Example 29. Figure 9.2 shows an example of a blockchain using Casper-FFG. The committed blocks include transactions t_1 to t_5 , where different stakeholders have deposited stake to obtain voting power. Note that:

- Not every block must include a transaction depositing stake.
- One block could also include multiple transaction depositing stake.
- Stake is usually deposited for a fixed time, e.g. only for the confirmation of the next 100 blocks.

Table 9.2 shows the stake deposited, that is relevant for voting on the three batched blocks shown in Figure 9.2. Note that t_6 is not included. That is because the block was not yet confirmed by the BFT algorithm.

Table 9.2 shows that node a holds $\frac{2}{7}$ of the voting power.




Figure 9.2: A blockchain showing transactions that deposit stake.

transaction	deposited stake	submitted by
t_1	2 eth	a
t_2	2 eth	b
t_3	1 eth	c
t_4	1 eth	d
t_5	1 eth	e
total	7 eth	

Table 9.2: Table showing voting power for Figure 9.2.

9.3 Algorand

Algorand is a hybrid blockchain that combines Proof of Stake with BFT.

Voting share is computed similar to Casper FFG. The main difference is that Algorand does not rely on Proof of Work for leader election. Algorand uses Verifiable random functions for leader election. Below we present a simplified variant:

1. Every block contains a *nonce*.
2. To check if a node is leader for block i he checks the following inequality:

$$H(\text{Sign}_{pk}(\text{nonce}_{i-1})) < \text{stake} \cdot d$$

Where d is a fixed difficulty, stake is a factor, adjusting the difficulty based on a nodes stake or voting power. $\text{Sign}_{pk}(\text{nonce}_{i-1})$ is a signature of the previous nonce, with a nodes public key.

3. If a node is a leader, he creates a new block, including

$$\text{nonce}_i = H(\text{Sign}_{pk}(\text{nonce}_{i-1}))$$

4. Instead of waiting for just one proposed block, nodes wait for time Δ and several proposed blocks. They then sign the proposed block with the smallest *nonce*.

To further improve scalability, Algorand uses the same mechanism used to select a leader, to select a committee, and requires only nodes in the committee to sign the certificate.

Chapter 10

Use Cases

This chapter is mainly based on a paper from Karl Wüst and Arthur Gervais titled Do you need a blockchain. The Figure 1 below is taken from the same paper.




Fig. 1: A flow chart to determine whether a blockchain is the appropriate technical solution to solve a problem (Table I should be considered in the decision making process as well). Writers refer to entities with write access to the database/blockchain, i.e. in a blockchain setting, a writer corresponds to a consensus participant. If a trusted third party (TTP) is available that is not always online, this can be used to establish a known group of writers, i.e. the TTP can function as a certificate authority in such a setting. Public and private permissioned blockchains differ in that a public blockchain allows anyone to read the contents of the chain and thus verify the validity of the stored data, while a private blockchain only allows a limited number of participants to read the chain. Note that for any blockchain based solution it is possible to make use of cryptographic primitives in order to hide privacy-relevant content.

10.1 Supply chain

Managing inventory, orders, processes, and failures in supply chains is a complex and expensive process. Improved data sharing between different parties has a big potential to improve supply chains. A blockchain may additionally improve traceability and prevent fraud.

Limitations:

- Data sharing is possible without a blockchain.
- Traceability is limited, due to the oracle problem, i.e., logged events on blockchain may not coincide with real world.
- Scalability might be an issue.

Opportunities: It seems especially for items that require thorough documentation, e.g., ivory or diamonds blockchain might be useful.

10.2 Inter bank payments

Today national inter bank payments are managed via national central banks. International inter bank payments are managed by multiple national banks. The system is expensive and complex.

A blockchain based approach may be a promising solution.

Limitations: Success depends on regulators. If appeals can be made by an external court, the system may become incorrect.

10.3 Digital services

When providing digital services, blockchains may have a good standing. Potential services are: Data storage, data exchange, dns resolution, search and lookup.

- Can systems be as trustful as centralized systems?
- Can systems improve privacy compared to centralized systems?