

Hyperledger Fabric


Execute-Order pipeline

Leander Jehl


Hyperledger fabric

- a toolbox to run your own, permissioned blockchain
- Permissioned:
 - Several well known participants are responsible to run the blockchain
 - Distrust each other.

Transaction processing pipeline




Transaction processing pipeline




- broadcast: send out transaction requires network resources
- validation: requires state
- ordering: requires coordination
- execution: requires state, must be deterministic.

Transaction processing pipeline




- broadcast: send out transaction requires network resources
- validation: access rights
- ordering: requires coordination
- execution: requires **What is the bottleneck?** deterministic.

Transaction processing pipeline



- For complex workloads, and small scale BFT systems, execution is the bottleneck.
 - Single threaded execution to be deterministic
 - Can be complex workloads
 - Execution has privacy concerns (need access to data)

Transaction execution




execution

Two approaches exist for crash fault tolerant systems:

- **Deterministic processing:** Each replica can process transaction and arrive at the same result.
- **Applying state change:** One replica executes transaction. Records state change Δ . All replicas apply Δ .


Transaction processing in Hyperledger fabric



- Execution happens before ordering.
- Execution policies, i.e. require n nodes to get the same result.
- Changes are submitted to ordering with signature from n nodes.
- During check, possibly inconsistent transactions are removed (aborted).

Transaction processing in Hyperledger fabric


State



- State is organized as (key, value) pairs.
- Execution result records, new values for certain keys and which keys have been read.
- Based on read and write keys, *check* can remove inconsistent transactions

Transaction processing in Hyperledger fabric

State



- State is organized as (key, value) pairs.
- Execution result records, new values for certain keys and which keys have been read.
- Based on read and write keys, *check* can remove inconsistent transactions. **For all stages but execution, values can be**

A BFT protocol

Simplified HotStuff

Leander Jehl

PoW or Certificates

Idea

PoW:

- In PoW the **difficulty decides the rate** at which the blocks are created, (throughput)
- High difficulty -> few forks
- Need to give time for a block to be propagated!
- But time to next block varies a lot!

PoW or Certificates

Idea


Goal:

- Throughput and confirmation only limited by time to propagate block, no extra waiting.

BFT protocol

Certificate vs. PoW

Idea: Send new block to nodes for validation and signature.
Then collect certificate.
new block




nodes

BFT protocol

Certificate vs. PoW


Idea: Send new block to nodes for validation and signature.
Then collect certificate.



BFT protocol

Certificate vs. PoW

Idea: Send new block to nodes for validation and signature.
Then collect certificate.



BFT protocol

Certificate vs. PoW

Idea: Send new block to nodes for validation and signature.
Then collect certificate.



$\langle \sigma_0, \sigma_1, \sigma_2 \rangle$



nodes

BFT protocol

Model

Model:

- We assume a permissioned system with $N = 3f + 1$ nodes.
- Nodes have unique ids and unique, known cryptographic keys.
- At most f of the nodes are byzantine faulty.



Certificate:

- A block has a certificate, if it contains signatures of $2f + 1$ nodes.

BFT protocol

Certificate vs. PoW



Idea: Send new block to nodes for validation and signature. Then collect certificate.



BFT protocol

Certificate vs. PoW



Idea: Send new block to nodes for validation and signature.
Then collect certificate.



BFT protocol

Certificate vs. PoW



Idea: Send new block to nodes for validation and signature.
Then collect certificate.



BFT protocol

Certificate vs. PoW

Idea: Send new block to nodes for validation and signature. Then collect certificate.



BFT protocol

Certificate vs. PoW

Idea: Send new block to nodes for validation and signature.
Then collect certificate.

Correct nodes sign only one block at given depth.

Obs: Faulty nodes may sign multiple blocks!


BFT protocol

Certificate vs. PoW problem

Idea: Send new block to nodes for validation and signature.
Correct nodes sign only one block at given depth.
Then collect certificate.

Problem: How to ensure that a certificate is created?

- Nodes may sign different blocks
 - No block gets a certificate
- **Solution:**



BFT protocol

Certificate vs. PoW problem

Idea: Send new block to nodes for validation and signature.
Correct nodes sign only one block at given depth.
Then collect certificate.

Problem: How to ensure that a certificate is created?

- Nodes may sign different blocks
- No block gets a certificate
- **Solution:** Leader

BFT protocol

Certificate vs. PoW problem

Idea: Send new block to nodes for validation and signature.
Correct nodes sign only one block at given depth.
Then collect certificate.

Problem: How to know that a certificate was created?

- A certificate may be collected by a single node
- The node with the certificate may fail and come back later
- **Solution:**

BFT protocol

Certificate vs. PoW problem

Idea: Send new block to nodes for validation and signature.
Correct nodes sign only one block at given depth.
Then collect certificate.

Problem: How to know that a certificate was created?

- A certificate may be collected by a single node
- The node with the certificate may fail and come back later
- **Solution:** Require multiple certificates

BFT protocol

Simple HotStuff (2 chain)

Preliminary:

- Every block includes a parent link (*previous block*).
=> Blocks form a **tree**.
Every block includes a **round number/view number**
- Every block must include a **certificate** for its parent
- A blocks' **round** must be larger than that of its parent

BFT protocol

Simple HotStuff (2 chain)

Rules

- *Every block must contain certificate for parent*
- *Every block must have round $>$ round of parent*
- **Rule 1:** After signing a block at round $max = r$,
a node may only sign at round $r' > max$.
- **Rule 2 preliminary:** After signing a block with parent p and only
sign blocks in subtree starting at p


Keep maximum value for max and $lock$ in local variables.




Simple HotStuff

Example

Example

- Nodes n_0 , n_1 , and n_2 sign block 






Node	<i>lock</i>
n_0	
n_1	faulty
n_2	
n_3	




Simple HotStuff

Example

Example

- Nodes n_0 , n_1 , and n_2 sign block 
- They set *lock* to 







Node	<i>lock</i>
n_0	
n_1	faulty
n_2	
n_3	




Simple HotStuff

Example

Example

- Nodes n_0 , n_1 , and n_2 sign block 
- They set *lock* to 
- n_3 signs 








Node	<i>lock</i>
n_0	
n_1	faulty
n_2	
n_3	




Simple HotStuff

Example

Example

- Nodes n_0 , n_1 , and n_2 sign block 
- They set *lock* to 
- n_3 signs 
- n_3 creates 









Node	<i>lock</i>
n_0	
n_1	faulty
n_2	
n_3	




Simple HotStuff

Example

Example

- Nodes n_0 , n_1 , and n_2 sign block 
- They set *lock* to 
- n_3 signs 
- n_3 creates 
- n_3 , n_1 , and n_2 sign block 




Node	<i>lock</i>
n_0	
n_1	faulty
n_2	
n_3	

Simple HotStuff

Example

Example

- n_0 sign block 




Node	<i>lock</i>
n_0	<input type="checkbox"/>
n_1	faulty
n_2	<input type="checkbox"/>
n_3	<input type="checkbox"/>




Simple HotStuff

Example

Example

- n_0 sign block 






Node	<i>lock</i>
n_0	
n_1	faulty
n_2	
n_3	




Simple HotStuff

Example

Example

- n_0 sign block 
- n_3 signs 






Node	<i>lock</i>
n_0	
n_1	faulty
n_2	
n_3	




Simple HotStuff

Example

Example

- n_0 sign block 
- n_3 signs 




Node	<i>lock</i>
n_0	
n_1	faulty
n_2	
n_3	




Simple HotStuff

Example

Example

- assume node n_0 is new leader.
- to get a certificate, n_0 must rely on faulty n_1



Node	<i>lock</i>
n_0	
n_1	faulty
n_2	
n_3	

BFT protocol

Simple HotStuff (2 chain)

Rules


- *Every block must contain certificate for parent*
- *Every block must have round $>$ round of parent*
- **Rule 1:** After signing a block at round $max = r$,
a node may only sign at round $r' > max$.
- **Rule 2:** After signing a block with parent p and $p . round = lock$,
only sign blocks with parents in round $pr \geq lock$


Keep maximum value for max and $lock$ in local variables.




Simple HotStuff

Example

Example

- assume node n_0 is new leader.
- to get a certificate, n_0 must
 - rely on faulty n_1
 - or extend 



Node	<i>lock</i>
n_0	
n_1	faulty
n_2	
n_3	

Solution:

n_0 waits for Δ time to get newest block

BFT protocol

Simple HotStuff

Def.:

a) A block with $round = r$ is **confirmed** if it has a child in $round = r + 1$, which has a certificate.

b) A block with $round = r$ is **confirmed** if it has a grandchild in $round = r + 2$

Theorem: *If a block is confirmed, only descendants of that block, can get a certificate.*

Proof: A majority of correct nodes have set their *lock* to the confirmed node.

BFT protocol

Simple HotStuff - Leader

Idea 1: Every round has a designated leader.

Idea 2: Nodes wait for Δ time for a proposal in current round, before accepting at next round.

BFT protocol

Simple HotStuff - Leader

Idea 1: Every round has designated leader.

Idea 2: Nodes wait for Δ time for a proposal in current round, before accepting at next round.

How can a leader avoid the situation from the example?

Ask all nodes for most recent certificate.

Wait for Δ time to receive proposal from all correct nodes.

BFT protocol

Simple HotStuff (3 chain)

Rules

- Every block must contain certificate for parent
- Every block must have round $>$ round of parent
- **Rule 1:** After signing a block at round $max = r$, a node may only sign at round $r' > max$.
- **Rule 2:** After signing a block with **grandparent** p and $p . round = lock$, only sign blocks with parents in round $pr \geq lock$

Keep maximum value for max and $lock$ in local variables.

BFT protocol

Simple HotStuff (3chain)

In this variant we need to wait longer for confirmation!

Def.:

- a) A block with $round = r$ is **confirmed** if it has a **grandchild** in $round = r + 2$, which has a certificate.
- b) A block with $round = r$ is **confirmed** if it has a **grand-grandchild** in $round = r + 3$

Theorem: *If a block is confirmed, only descendants of that block, can get a certificate.*

Proof: A majority of correct nodes have set their *lock* to the confirmed node.

BFT protocol

Simple HotStuff - Leader

Idea 1: Every round has designated leader.

Idea 2: Nodes wait for Δ time for a proposal in current round, before accepting at next round.

How can a leader avoid the situation from the example?

Ask all nodes for most recent certificate.

Wait for $2f + 1$ replies

No leader needs not wait for Δ time!